

POLITE LANGUAGE REFLECTS COMPETING SOCIAL GOALS

A DISSERTATION
SUBMITTED TO THE DEPARTMENT OF PSYCHOLOGY
AND THE COMMITTEE ON GRADUATE STUDIES
OF STANFORD UNIVERSITY
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

Erica J. Yoon

April 2019

© Copyright by Erica J. Yoon 2019
All Rights Reserved

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Michael C. Frank) Principal Adviser

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Ellen Markman)

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Noah Goodman)

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Hyowon Gweon)

Approved for the Stanford University Committee on Graduate Studies

Abstract

We use polite speech on a daily basis. From “thanks” and “please” to “your dress is cute” to “these cookies could use a bit of salt,” people often produce polite utterances that are indirect or even false to some degree. Why do people speak politely? This thesis proposes a goal-based account of polite speech, that polite speech arises from a set of competing social goals: the speaker’s desire to transfer information in the most truthful and informative manner possible (“informational goal”), and to abide by social norms and expectations and/or maintain the interactants’ *face* or public self-image (“pro-social goal”) and/or present herself as a particular kind of individual (e.g., kind, helpful person; “self-presentational goal”). In Chapter 1, I provide an overview of this integrative theoretical framework that aims to unify previous theoretical accounts of polite speech and explain existing empirical studies on understanding and production of polite speech in adults and children. In Chapter 2, I provide a computational model that formalizes the notion of goals as utilities that speakers try to maximize through language use, and show that this model successfully captures adults’ predictions and judgments for polite lies and indirect speech. Then I present two sets of empirical studies looking at the development of polite language understanding: 2- to 4-year-old children’s judgments for polite requests (Chapter 3) and 5- to 8-year-old children’s judgments for polite lies versus blunt truths (Chapter 4). Overall, the work presented in this thesis reveals how adults and children’s understanding of polite speech reflects their understanding of speaker goals to be informative and social.

Dedication

FIXME

Acknowledgments

FIXME

Contents

Abstract	iv
Dedication	v
Acknowledgments	vi
Introduction	1
1 A goal-based account of polite language	5
1.1 Introduction	5
1.2 Part I: A Goal-based account of politeness	7
1.2.1 Informational goal	8
1.2.2 Prosocial goal	11
1.2.3 Presentational goal	12
1.3 Part II: Previous theoretical accounts of polite speech	13
1.3.1 Politeness as observance of communicative maxims	15
1.3.2 Politeness as strategy for facework	16
1.3.3 Politeness as social rules and norms	18
1.3.4 Politeness in the game-theoretic approach	20
1.4 Part 3: Empirical work on polite speech in adults	22
1.4.1 Interpretation of polite speech	22
1.4.2 Evaluation of polite speakers' intentions	23
1.4.3 Production of polite speech	24
1.5 Part 4: Empirical work on polite speech in children	26

1.5.1	Rule-based polite speech	26
1.5.2	Indirect speech	28
1.5.3	White lies	29
1.5.4	Summary	31
1.6	Conclusions	31
2	Adults consider tradeoffs between competing social goals to predict polite language use	33
2.1	Introduction	33
2.2	Model predictions	38
2.3	Experiment: Speaker production task	40
2.3.1	Participants	41
2.3.2	Design and Methods	41
2.3.3	Behavioral results	41
2.3.4	Model results	42
2.4	Discussion	44
3	Children understand social goals behind polite requests	47
3.1	Introduction	47
3.2	Experiment 1	50
3.2.1	Methods	51
3.2.2	Results and Discussion	52
3.3	Experiment 2	55
3.3.1	Methods	55
3.3.2	Results and Discussion	56
3.4	Experiment 3	57
3.4.1	Methods	57
3.4.2	Results and Discussion	58
3.5	General Discussion	59
4	Children understand tradeoffs between informational and prosocial goals	61

Conclusion	62
A Supplementary materials for Chapter 2	68
A.1 Model details	68
A.2 Literal semantic task	69
A.2.1 Participants	69
A.2.2 Design and Methods	69
A.2.3 Behavioral results	69
A.3 Data analysis	70
A.4 Full statistics on human data	71
A.5 Model fitting and inferred parameters	71
A.6 Data Availability	72
A.7 Supplemental Figures	72
B Supplementary materials for Chapter 3	76
B.1 Model output for Experiment 3.1	76
B.2 Model output for Experiment 3.2	76
References	77

List of Tables

1.1	Review of existing theories of polite language	14
2.1	Comparison of variance explained for each model variant and log Bayes Factors quantifying evidence in favor of alternative model in comparison.	43
2.2	Inferred phi parameters from all model variants with more than one utility.	44
A.1	Predictor mean estimates with standard deviation and 95% credible interval information for a Bayesian linear mixed-effects model predicting negation production based on true state and speaker goal (with both-goal as the reference level).	71
A.2	Inferred negation cost and speaker optimality parameters for all model variants. . . .	72

List of Figures

1	Schematic overview of the dissertation content.	4
2.1	Diagram of the model: The polite speaker observes the true state and determines her goal between three utilities (informational, social, and presentational), and produces an utterance.	36
2.2	Model behavior. Listener inferences about the true state (e.g., the rating truly deserved by the poem; A) and the speaker's utility weighting (ϕ_{S_1} or how informational vs. social the speaker is, where $\phi_{S_1} = 0$ is fully social, and $\phi_{S_1} = 1$ is fully informational; B) as a function of the utterance heard (facets). C: Purely self-presentational speaker production behavior as a function of the kind of speaker they wish to present themselves as (facets; relatively more informational, e.g., $\phi_{S_1} = 0.05$, vs. social as represented, e.g., $\phi_{S_1} = 0.95$).	39
2.3	Example of a trial in the speaker production task.	40
2.4	Full distribution of human responses vs. model predictions. Error bars represent 95% confidence intervals for the data (vertical) and 95% highest density intervals for the model (horizontal).	42
2.5	Comparison of predictions for proportion of utterances chosen by pragmatic speaker from possible model variants (left) and human data (rightmost) for average proportion of negation produced among all utterances, given true state of 0 heart (on a scale of 0 to 3) and speaker with both goals to be informative and kind. Gray dotted line indicates chance level at 12.5%.	43

3.1	Bottom right: Story example. Top, left: Results. Proportion of correct responses to questions comparing between a speaker who used a politeness marker (where blue indicates "please", yellow "can you", and red "can you please") versus a speaker who did not. Data are binned into one-year age groups. Each row represents data from a different Experiment. Columns represent different questions asked. Dashed line represents chance level (i.e., if participant were guessing at random).	53
A.1	Semantic measurement results. Proportion of acceptances of utterance types (shown in different colors) combined with target words (shown in different facets) given the true state represented on a scale of hearts. Error bars represent 95% confidence intervals.	70
A.2	Experimental results (solid lines) and fitted predictions from the full model (dashed lines) for speaker production. Proportion of utterances chosen (utterance type "direct vs. indirect" in different colors and words shown on x-axis) given the true states (columns) and speaker goals (rows). Error bars represent 95% confidence intervals for the data and 95% highest density intervals for the model. Black dotted line represents the chance level.	73
A.3	Comparison of predictions for proportion of utterances chosen by pragmatic speaker from possible model variants (left) and human data (rightmost) for average proportion of negation produced among all utterances, given true state of 0 heart and speaker with a goal to be informative (top), kind (middle), or both (bottom). Gray dotted line indicates chance level at 12.5%.	74
A.4	Experimental results (left) and fitted model predictions (right) for average proportion of negation produced among all utterances, given true states (x-axis) and goals (colors).	75

Introduction

We use and hear polite speech on a daily basis, ranging from simple words of apology (“sorry”) or gratitude (“thanks”) to compliments (“I love your dress!”) and requests (“Can you please open the window?”). Adults and even young children spontaneously produce requests in polite forms (Axia & Baroni, 1985; Clark & Schunk, 1980). Speakers exhibit politeness strategies even while arguing, preventing unnecessary offense to their interactants (T. Holtgraves, 1997). Listeners even attribute ambiguous speech to a polite desire to hide a truth that could hurt another’s self-image (e.g., Bonnefon, Feeney, & Villejoubert, 2009). In fact, it is difficult to imagine human speech that efficiently conveys only the truth. Intuitively, politeness is one prominent characteristic that differentiates human speech from stereotyped robotic communication, which may try to follow rules to say “please” or “thank you” yet still lack genuine politeness.

Although language users use polite speech on a daily basis, explaining why we use polite speech or how we understand it is not as straightforward as it first seems. While very simple polite utterances can be produced from straightforward rules (e.g., say “sorry” when you did something bad to someone), When speakers want to tell the listener to “close the window,” they often use a more roundabout way and say “can you please close the window?” When people see that their interactant is wearing a new outfit that they think is hideous, they might still say “Your dress looks gorgeous!” As such, polite utterances often seem to misrepresent their intended message or conceal the truth, which shows that polite speech violates a critical principle of cooperative communication: exchanging information efficiently and accurately (Grice, 1975).

If politeness only gets in the way of effective information transfer, why be polite? Clearly, there are social concerns, and most linguistic theories assume utterance choices are motivated by these concerns, couched as either polite maxims (Leech, 1983), social norms (Ide, 1989), or aspects of

a speaker and/or listener’s identity, known as *face* (P. Brown & Levinson, 1987; Goffman, 1967). All of these theories use different approaches to explain polite language, and some are even framed as counterarguments to existing theories (e.g., see Richard J Watts (2003) and Matsumoto (1988) responding to some issues in P. Brown & Levinson (1987)). One possible commonality among these theories however, is that they all describe ways in which language communication deviates from certain expected utterances or conversations due to speakers’ social concerns.

In this thesis, my goal is to offer an integrative theoretical framework that aims to unify these existing theories, and provide empirical evidence in support of this framework. Specifically, I argue for a *goal-based* theory of polite speech: that polite utterances arise from competing social goals that speakers have, such as their desires to convey information as truthfully and efficiently as possible (“informational goal”), to make the listeners feel happy and respected and thereby boost or maintain their face (“prosocial goal”), and to present speakers themselves in a good light (e.g., that they are kind and helpful; “presentational goal”). Speakers then have to consider the tradeoff between these goals, and think about which goal to prioritize and how much to do so to determine their utterance.

For example, imagine that Alice and Bob are having a conversation and Bob asks for Alice’s feedback on his cookies that he baked (“How did you like my cookies?”) and Alice thinks the cookies tasted bad and salty (Figure 1, top panel). Alice’s utterance would differ depending on her goals: whether she wants to prioritize informational goal or telling the truth to Bob; social goal or making Bob feel happy; or presentational goal or presenting Alice herself in a good light that she is kind (Predictions of this specific scenario will be explained in detail in Chapter 2).

The contents of this dissertation will be as follows, as shown in Figure 1: In Chapter 1 (top panel of Figure 1) I present an integrative goal-based framework that aims to explain polite speech based on the idea that it reflects a tradeoff between competing social goals that speakers have. Then using this framework, I will explain existing empirical studies on understanding and production of polite speech in adults and children. Chapters 2-4 describe a set of computational and empirical studies of children and adult’s understanding of polite language (bottom panels of Figure 1) . In Chapter 2, I provide a computational model that formalizes the notion of goals as utilities that speakers try to maximize through language use, and show that this model successfully captures adults’ predictions and judgments for polite lies and indirect speech. Then I present two sets of empirical studies looking at the development of polite language understanding: Chapter 3 examines 2- to 4-year-old children’s

judgments for polite requests, and Chapter 4 looks at 5- to 8-year-old children's judgments for polite lies versus blunt truths.




Figure 1: The upper panel shows a schematic overview of an integrative framework of polite language understanding based on competing social goals. The lower panels show different studies examining adults' and children's understanding of different component goals (and possible tradeoffs between them) that correspond to each chapter of the dissertation.

Chapter 1

A goal-based account of polite language

1.1 Introduction

Imagine that a stranger on the street approached you and asked: “I’m sorry to bother you, but could you tell me the way to the city hall?” Regardless of your answer, you probably would not feel puzzled or offended by the way in which the stranger decided to seek information that he needs. This is in contrast with a different situation where the stranger said to you instead: “Tell me the way to the city hall.” In such case you would immediately notice the lack of politeness in his utterance, and your response to him may be negatively affected by the irritating oddity of the situation. Now imagine another context where a person was wearing a new, flashy dress, and her friend thought the dress was hideous. It may actually be more surprising if the friend truthfully said “Your new dress is ugly,” than if the friend decided to lie and say “Your new dress is gorgeous!” But why? Why are people expected to speak politely, when there are alternative utterances that can convey information about the world or the speaker’s intention more directly (“Tell me the way”) or truthfully (“Your dress is ugly”)?

Language is a virtuous tool that serves many functions. Through language, people communicate information about the world, but also form their social relationships and establish their identity within their society. On one hand, some theories of language functions describe language as a

transmission device for transferring from a sender to a receiver information that reflects context or the state of affairs (B  ijhler, 1934; Jakobson, 1960; Shannon, 1948). The importance of informativity is further emphasized in more recent, influential theories on pragmatics of natural language, which explain how meanings beyond literal meanings of utterances arise (Grice, 1975; Searle, 1975). On the other hand, some linguistic theories, especially those with references to language development, identify social roles of language that people use to make contact with others and form relationships (Ervin-Tripp, 1967; Halliday, 1975). These theories underscore how linguistic rules that language users tend to follow represent the norms and structure of the community using the language (Ervin-Tripp, 1969).

Could polite speech reflect both the informational and the social roles of linguistic communication? Previous theoretical accounts of polite speech vary in their focus on informational versus social aspects of polite language. Some theories view polite speech as reflecting social rules and norms (Richard J. Watts, Ide, & Ehlich, 1992), some as abiding by communicative maxims that people are expected to follow in conversations, to be both as informative and as affirmative of their conversational partner as possible (Lakoff, 1973; Leech, 1983), and yet some others as performing face management, or trying to maintain interactants' good public self-image or reputation (P. Brown & Levinson, 1987).

In this Chapter, I propose that that polite speech highlights both informational and social uses of language: Polite speech reflects a principled tradeoff between the informational, epistemic content a speaker wants to convey (e.g. "I want you to tell me the way to the city hall") and other social concerns, such as prosocial or self-presentational goal that the speaker wants to express for herself and others ("I'm not rudely commanding you to tell me the answer, but making a request in a respectful way"). Thus, my goal is to unify previous theoretical frameworks in one, goal-driven account of polite speech. In what follows, I will describe the goal-based account of polite speech in detail (Part I), and summarize previous models and theories of polite speech and situate them within the framework of the current goal-driven account (Part II). Then I will examine empirical evidence for goal-driven approach to polite speech: In Part III, I will focus on empirical work on adult production and comprehension of polite speech which show that adults reason about speaker goals in polite speech; and in Part IV, I will probe empirical evidence from developmental work that children's production and understanding of polite speech advance as they grow older. In doing this

I will show that children’s production and comprehension of polite speech is related to the relative complexity of polite speech based on its goal tradeoff.

1.2 Part I: A Goal-based account of politeness

What does it mean to speak politely? Common instances of polite speech that occur to one’s mind probably include the simplest politeness markers, such as “please,” “thanks,” and “sorry.” More complicated examples would involve ways in which, for example, a person make a request: under normally conceivable circumstances, it would certainly be more polite to ask “Would it be too much trouble to ask you to complete this survey when you’re not too busy?” than to say “Do this survey now.” The word “polite” can sometimes carry a negative undertone in its meaning, as in “she was just being polite,” which is likely to mean the speaker was hiding her genuine beliefs or intentions to make the listener feel good. From these examples, we can identify a few characteristics that a polite utterance may exhibit: observance of social rules, relatively high degree of elaborateness or indirectness, and dishonesty or disingenuousness in the interest of others’ feelings or reputations.

More formal definitions of the term politeness also reveal common features that polite speech shows. Cambridge and Oxford Dictionaries respectively define what it means to be polite: “behaving in a way that is socially correct and shows respect for other people’s feelings”; and “courteous, behaving in a manner that is respectful or considerate of others; well-mannered” (“Polite,” 2017a, 2017b). Similar to the previous examples of polite speech, these definitions suggest politeness involves (1) observance of social expectations and (2) respect for others. Boyer (1702)’s *The English Theophrastus: of the Manners of the Age*, compilation of texts describing the English life in the early eighteenth century, identifies a purpose in trying to be polite: “Politeness may be defined as a dextrous management of our Words and Actions whereby *men make other people have a better Opinion of us and themselves* [emphasis added].” Thus, according to the Theophrastus definition, speakers speak politely in order to boost the self-images of the interactants (both the speakers themselves and their addressees).

These common themes of politeness have been identified by previous theoretical accounts of polite speech (reviewed in detail in Part II). But each of the accounts only focused on certain aspects of polite speech but disregarded others, and their explanations for politeness have been viewed as largely disparate. Here I make a unification proposal, where these existing approaches to polite

speech can be united under a single goal-based account of polite speech.

I propose that polite speech reflects some degree of tradeoff between three main communicative goals: informational, prosocial, and (self-)presentational. *Informational goal* has to do with the speaker’s desire to convey the most accurate information in the most efficient manner. *Prosocial goal* is about the speaker’s desire to retain the listener’s acceptable self-image as a decent individual and as a reputable member of society. *Presentational goal* reflects the speaker’s desire to present the speaker herself in a good light, to appear to be a kind and helpful individual. Below I describe each goal in more detail.

1.2.1 Informational goal

A speaker’s informational goal, i.e. to prioritize information transfer in communication, may involve two closely related notions: informativity and truthfulness. Informativity is the notion of conveying the intended meaning in the most efficient and precise manner possible. The idea of informativity here is similar to Grice (1975)’s notion of cooperativity: A speaker will cooperatively choose utterances such that the listener can understand her intended message. Thus, the current notion of informativity that I adopt will encompass the whole Cooperative Principle (CP) that Grice posited (“Make your contribution such as required by the purposes of the conversation at the moment”), and especially the Maxim of Quantity (“Make your contribution as informative as is required”), though Maxims of Relevance (“Be relevant”) and Manner (“Be perspicuous”; i.e. be brief, orderly, and unambiguous) can also be relevant to the notion of informativity as I discuss in this Chapter.

The idea of informativity has been formalized in probabilistic (Bayesian) models as a utility function of a speaker with particular goals in mind. The “rational speech act” (RSA) theory of language understanding (see N. D. Goodman & Frank (2016a) for a review) assumes that listeners expect speakers to aim to be helpful yet parsimonious, choosing their utterances approximately optimally based on a communicative goal (e.g., inform the listener) and interpret an utterance by inferring what the helpful speaker meant based on the utterance and any other relevant information about the world. The theory defines a standard, informative utility as the amount of information a literal listener (L_0) would still not know about world state (s) after hearing a speaker’s utterance (w):

$$U_{epistemic}(w; s) = \ln(P_{L_0}(s|w))$$

where the utterance choice is approximately rational (i.e., in proportion to the expected utility gain) and w is chosen from a set of alternative, relevant utterances.

For example, if Bob asked Alice “How was my cookie?” and Alice said to Bob “It was good,” with only the goal to be informative in mind, Bob would think that Alice was being maximally informative by using the word “good” instead of another relevant, stronger word such as “amazing,” and infer that Ann meant “good but not amazing” because otherwise Ann would have used the word “amazing” instead.

I note here that Ann’s speech act could be analyzed as having observed the Gricean Maxim of Quantity, by making her utterance maximally informative, and Bob’s inference as being based on such assumption that Ann’s utterance is as informative as is required to meet Bob’s needs. But as the comparison between the former goal-directed analysis and this latter maxim-based analysis may reveal, the maxim-based account is difficult to formalize (Hirschberg, 1985) whereas the goal-directed view allows for quantitative account of factors contributing to the linguistic phenomena at hand (N. D. Goodman & Frank, 2016a). From here on, I take the goal-directed view rather than the Gricean maxim-based view of the speaker; thus, analyses of speakers and their utterances will be based on speakers’ communicative goals to convey information, etc., rather than their observation or violation of Gricean maxims.

Informational goal can also involve truthfulness: the meaning that accessed by the listener should match the true state of the world as closely as possible. In other words, the speaker will want to convey what is true (to the extent of her knowledge), not what is false. For example, if Bob baked some cookies and they tasted terrible, and Alice had the goal to be truthful, she would want to convey what matches the true state of the world as closely as possible (“Your cookies tasted terrible”). On the other hand, if Alice remarked to Bob about his cookies “Your cookies tasted good,” with goal to be informative and truthful, the true state of the world must be such that Bob’s presentation was truly good (but probably not amazing, because otherwise Alice would have said that it was amazing), and not bad or terrible.

As for when speakers decide not to tell the truth, there is a difference between violating versus flouting of truthfulness (which Grice also discusses; see Grice (1975), p. 49, 53). Flouting involves

contradicting common knowledge shared between the speaker and listener about the true state of the world, such that the listener notices that the utterance meaning does not match the true state. For example, after Alice and Bob together watch a movie that was obviously gory and disturbing to both of them, if Bob says “well, that was a really happy, fluffy movie!” then his utterance would be flouting truthfulness goal, and Ann would recognize Bob’s utterance as an ironic one. On the other hand, violating truthfulness does not hold assumption of such common knowledge of the true state of the world, and thus the listener may not notice the mismatch between the utterance meaning and true state of the world, although the listener can potentially access it through other means (e.g. realizing that the speaker had reasons to lie). Here I will mainly focus on violation, not flouting, of truthfulness, for example when speakers tell white lies (i.e., when the speakers do not intend that the listeners know what the truth is).

Speakers’ informational goal to be informative and truthful may encompass communicative co-operation in both locutionary and perlocutionary senses. A locutionary goal deals with conveying the intended meaning of the utterance within a conversation, whereas a perlocutionary goal involves achieving the speaker’s ultimate goals toward the listener (Attardo, 1997). What could be an informational goal in its perlocutionary sense? In being truthful, speakers may ultimately want to maintain their moral obligation to tell the truth to others. This obligation is in line with Western philosophers’ argument throughout the history that it is morally wrong to lie (Augustine, 1952; Kant, 1949), although there have been debates on whether the degree of wrongness may depend on context (e.g., if the speaker was telling a white lie; Sweetser, 1987). For example, if Alice said to Bob “Your cookies were good,” Alice’s locutionary goal would be to convey to Bob her intended meaning that his presentation was good (but perhaps not amazing), whereas her perlocutionary goal would be for Bob to think that the presentation was good, which was (apparently) the truth; Alice thereby upholds her moral obligation to tell the truth to Bob. Alice’s goal to be truthful, then, is achieved in both locutionary and perlocutionary senses.

Besides an informational goal, a speaker may also want to address concerns that are social in nature: having to do with interacting and maintaining good relationships with other people. Below I describe two related but different goals that speakers may want to accomplish for social reasons: prosocial and presentational.

1.2.2 Prosocial goal

A prosocial goal involves the speaker's desire to follow social norms and make others feel happy and respected. Speakers can try to accomplish the prosocial goal in several ways, one of which is social norm observance: abiding by social norms and expectations. There may be simple rules such as "say please when you make a request" or "say thanks to express gratitude," but sometimes the norms can be more complex. Speakers should avoid saying utterances that are *too* polite, to the extent that the utterances become marked and are no longer considered "optimally polite." For example, a request for opening a window by saying "Sorry, could you open that window behind you? Thanks." would be a normal, socially expected way to make the request; but a request such as "I'm so terribly sorry to bother you with this irritating request, but if you don't mind, would you care to open that window behind you, only if it's not going to be too much trouble for you?" would be a signal to the addressee that something in the situation is odd and marked; either that the request involves a higher cost than is normally expected for opening a window, or the speaker is unusually afraid of incurring a debt to the addressee, etc. This principle of social norm observance is then parallel to Grice's Cooperative Principle, in that the CP outlines normative expectations for a speaker who wants to convey information as efficiently as possible, whereas the current principle of social norm observance deals with normative expectations for a speaker who wants to maintain social order. Thus, if the CP is a principle of information transfer, social norm observance is a principle of social order. Both principles call for unmarkedness of utterances, and when the utterances are marked due to a violation of its rules, then the listener will try to infer reasons for such violation.

Speakers may also try to be prosocial through face management. *Face* is a notion introduced by Goffman (1967), and represents an individual's publicly manifest self-esteem. He argued that people perform interpersonal rituals whereby face maintenance is a fundamental condition of the interactions. Goffman identified two kinds of faces that people want to maintain: *positive face*, or the want for solidarity or approval from others, and *negative face*, or the want to be free from imposition. Interactants will always want to preserve each other's face, and so potential face threats will somehow have to be modified. P. Brown & Levinson (1987) suggested that a strategy for such facework is politeness, which they defined as deviation from Gricean informativity (described in detail in Part II).

For example, a request such as "You couldn't possibly pass the salt, could you?" would be an

example of negative politeness strategy (i.e. a strategy to save negative face; P. Brown & Levinson, 1987, p. 136) as the speaker is being pessimistic about the compliance of her request and not assuming that the listener has to be willing or able to do any acts predicated of him. On the other hand, utterances that emphasize the common ground between the speaker and listener (i.e. that the speaker and listener share the same goals, values, knowledge, etc.), and address the fulfillment of the listener's want are positive politeness strategies; for example, "What a beautiful vase this is! Where did it come from?" (P. Brown & Levinson, 1987, p. 103) saves the listener's positive face by attending to the listener's wants and interests. When face management is in conflict with informational goals, the meaning of utterance would differ depending on which goal the speaker decided to prioritize. As described earlier, if Ann said to Bob, "Your presentation was good," and she wanted to prioritize informational goals only, then her utterance would indicate that Bob's presentation was truly good (though perhaps not amazing). However, if Ann spoke with a prosocial goal to save Bob's face and wanted to boost his self-image instead, then Bob's presentation actually could be bad rather than good.

1.2.3 Presentational goal

Language also reflects a speaker's goal to present themselves in a good light, thereby saving the speaker's own face. This last goal is related to the informational and prosocial goals previously described, in that speakers must be mindful of the listener's want to be informed or to maintain his positive self-image, but instead of actually *being* maximally informative or prosocial, presentational goal concerns *appearing* to care about these goals. Thus, a speaker may engage in a recursive reasoning about a listener who thinks about a speaker who wants to be informative and/or prosocial, and then can produce utterances to make the listener *think* that the speaker is being informative, being prosocial, or both of those things. For example, rather than saying "your talk was terrible," people are more likely to say "it wasn't bad" to *indirectly* suggest that the talk was not great, while signaling their good intention to be nice and not say the harsh truth (see Chapter II for the formal definition and more detailed description of the presentational goal).

1.3 Part II: Previous theoretical accounts of polite speech

In Part II, I aim to (i) describe different classes of theoretical approaches to the understanding of polite speech; (ii) for each class, explain how the approach can be situated within the current proposal for the goal-based account for polite speech; and (iii) discuss what advantages the goal-based account can offer beyond the existing approaches. Summary of prominent theories and their implications within goal-based framework can be found in Table 1.1.

Table 1.1: Summaries of previous theoretical approaches to polite speech and their implications within the current goal-based framework.

Politeness as:	Theories offered by:	Summary	Advantage of goal-based framework
Observance of communicative maxims Strategy for facework	Lakoff (1973); Leech (1983)	Polite speech is governed by conversational rules and principles, that are complementary to Gricean Cooperative Principle (Grice, 1975)	Gradient degree of goal tradeoff can be represented, instead of binary observance of different maxims
	Brown & Levinson (1987)	Speakers produce polite speech to save interactants' face (positive and negative)	Clearer distinction between notions of truthfulness and informativity becomes possible, which then allows more precise analysis of goal tradeoffs
	Spencer-Oatey (2000)	Speakers try to preserve interactants' face and sociality rights	
Social rules and norms	Watts (1989); Locher & Watts (2005); Watts et al. (1992); Watts (2003)	Speakers use speech to perform relational work (not only face-work); speakers produce im/polite (marked) utterances and politic (unmarked, normative) utterances	No distinction between polite vs. politic utterances is necessary
	Fraser & Nolen (1981); Matsumoto (1988); Ide (1989); Mao (1994)	Speakers want to fulfill societal obligations	Facework and social obligation do not have to be mutually exclusive
Model of game theory	Franke & Jager (2016); Pinker et al. (2008); Van Rooy (2003); Quinley (2011)	Speakers use polite speech to get what they want while allowing for plausible deniability, or to communicate their good intentions (while incurring cost)	Self-interest is formalized as arising from recursive reasoning that is based on genuine other-oriented goals

1.3.1 Politeness as observance of communicative maxims

One approach to polite speech is largely based on the Gricean perspective, and claims that conversation is driven by general communicative principles, and politeness is a maxim (or made up of maxims) that accompanies other principles. Whereas Grice focused on principles that speakers follow to make their speech maximally efficient for information transfer, he also noted: “There are . . . all sorts of other maxims . . . , such as ‘Be polite,’ that are normally observed by participants in talk exchanges” (Grice, 1975, p. 47). Searle (1975) discussed conditions of performing indirect speech acts (e.g. “Can you pass the salt?”) for which, he claimed, “politeness is the chief motivation” (p. 76). Thus, Grice and Searle both identified politeness as a driving factor in communication, but the concept of politeness was largely undeveloped at the time.

The communicative maxim approach for polite speech expanded with proposals for conversational rules and principles that govern polite speech. Lakoff (1973) suggested a set of rules (“Don’t impose,” “Give options,” “Be friendly”) that underlie utterances that are polite, and thus deviate from directly expressed meanings. Leech (1983) developed a similar proposal to Lakoff’s in greater detail, proposing Politeness Principle (PP), which is complementary to the Gricean CP. Leech argued that the PP, like the CP, can be subclassified into more specific maxims: Maxims of Tact, Generosity, Approbation, Modesty, Agreement, and Sympathy. The primary postulate of the PP is that interactants prefer to express polite beliefs, which are beliefs that are favorable to the other person (and/or unfavorable to oneself). For example, Approbation Maxim states that a speaker observing the PP will minimize dispraise, and maximize praise, of the other person (e.g. “Your dress looks gorgeous!”), whereas Modesty Maxim states that a speaker will minimize praise, and maximize dispraise, to the speaker herself (e.g. “This is just a small gift, but I hope you like it” downplaying the value of the gift). Gu (1990) proposes addition of “Balance Principle” to the set of politeness maxims, where favors from A to B are balanced by favors from B to A, such that the PP can function to maintain social equilibrium.

The theories of politeness as a communicative principle focuses on cases of tradeoff between speakers’ communicative goals: On one hand, the speaker wants to be informative, following Gricean maxims, but on the other hand, the speaker wants to follow social rules, represented by Lakoff’s rules or Leech’s PP. An utterance then reflects some level of both of these desires. However, there arise a classification and quantification issue: maxims of PP are informal and categorical, which makes

it difficult to represent the degree to which the utterance observes or violates a given maxim. In this regard, the goal tradeoff approach is preferable to the maxim approach, for it becomes possible to represent a gradient degree of tradeoff between speakers' goals of informativity and social rule following.

1.3.2 Politeness as strategy for facework

Another approach to analysis of polite speech that has been highly influential is the model of politeness as face management, developed by P. Brown & Levinson (1987). The model is primarily based on the concept of face (Goffman, 1967) that deals with an individual's public persona (introduced earlier in Part 1). P. Brown & Levinson (1987) take a strategic approach to politeness, and focuses on strategies that speakers employ in order to avoid, redress, or mitigate threats to face (either the addressee's or the speaker's own). Like the communicative maxim approach, P. Brown & Levinson (1987)'s account also start at the assumption of the Gricean CP, and attribute the cause of a speaker's deviations from the CP to the speaker's desire to be polite. Thus, given a speaker's desire to save face, the more an utterance deviates from maximal informativity, the more "polite" (face-saving) the utterance will be.

Within the goal-based framework, it can be said that P. Brown & Levinson (1987) recast the notion of a cooperative speaker as one who has both an informational goal to improve the listener's knowledge state as well as a social goal to minimize any potential damage to the hearer's (and the speaker's own) face. With the idea of these conflicting goals, Brown and Levinson's theory is conceptually closest to the current proposal among all theories to be discussed. Specifically, they primarily focus on the conflict between goals of social face management and epistemic informativity. For example, by saying "Can you please open the window?" instead of "Close the window," the speaker decides to sacrifice informativity (i.e. maximally efficient transfer of the intention for the addressee to open the window) in order to save the addressee's negative face (i.e. freedom from imposition or order from others). Another example of negative politeness (i.e. strategy to save negative face) is a speaker fronting his gift-giving by saying "This is just a small gift" in order to emphasize that the listener is not incurring too much debt to the speaker and thus not being imposed a burden to return the favor. On the other hand, positive politeness or act of saving positive face can be exemplified by utterances that emphasize approval of the listener's interest or performance,

e.g. “What a fantastic garden you have!” (P. Brown & Levinson, 1987, p. 104).

One drawback in P. Brown & Levinson (1987) is that key elements required for analysis of speakers’ goals or intended meanings are ambiguous and difficult to formalize. The core assumption of Brown and Levinson’s analysis is that deviations from Gricean maxims (that prioritize the match between the literal and intended meanings) lead to increase in politeness. In order to estimate the degree to which speakers try to be polite, there should be a way to measure the degree of deviation from the Gricean maxims (i.e. the degree of mismatch between the utterance surface meaning, intended meaning, and true state of the world; see Section 1 for how these factors can be formalized e.g. in RSA). However, Brown and Levinson do not provide a way to formalize the literal or intended meaning or the state of the world. For example, it is unclear how much deviation from the Gricean maxim(s) occurred when a speaker produced an utterance such as the example previously mentioned: “What a fantastic garden you have!” Since the authors do not provide the true state of the world (i.e. how beautiful the garden actually was in the speaker’s opinion), it is difficult to define or quantify the speaker’s effort for politeness or the cost of face threat the message would have involved given an alternative utterance (“It is a mediocre garden you have.”) A clearer identification of the true state of the world, the speaker intended meaning, and the set of relevant alternative utterances, and a way to formalize these elements, will help with quantificational analysis of polite utterances.

Spencer-Oatey (2000) put forward another theory based on Goffman’s notion of face; it is a more general model of rapport management, or management of interpersonal relations, distinguished from face management as proposed by Brown and Levinson. Spencer-Oatey challenges Brown and Levinson’s distinction between positive and negative face, and argues that the former involves the concept of face, or the positive social value claimed by a person, whereas the latter does not concern face but rather what she calls “sociality rights,” or a person’s entitlements in interactions with others. She then further proposes that face management and sociality rights management each has both personal components (concerning self-esteem and individual identity) and social components (concerning social role and entitlements within relations with others). Spencer-Oatey then largely focuses on speakers’ attempt to balance between general epistemic goals and both face management and social expectations. Thus, in saying “This is just a small gift, but I hope you like it” the speaker may not only be concerned about the speaker’s and the listener’s faces as individuals but also the

conventional social obligations associated with gift-giving that the speaker and listener are expected to follow through. However, similar to Brown and Levinson's analysis, Spencer-Oatey's analysis is limited in that it does not identify or formalize the epistemic knowledge (the speaker intended meaning and true state of the world) behind utterances, making it difficult to quantitatively analyze the goal tradeoffs.

1.3.3 Politeness as social rules and norms

Another set of theories for polite speech relies on the notion that politeness deals with following social norms and expectations. One line of theory argues for the need to distinguish between marked, strategic polite behaviors versus unmarked, normative "politic" behaviors. Richard J Watts (1989) defines "politic" behavior as "socio-culturally determined behavior directed towards the goal of establishing and/or maintaining in a state of equilibrium the personal relationships between the individuals of a social group, whether open or closed, during the ongoing process of interaction." For example, "May I open the window behind you?" will be a normative, politic statement if used in a formal setting or said to a stranger, but can be considered marked, overly polite statement when said to a close family member.

Richard J. Watts et al. (1992) posited that examples of polite speech in Brown and Levinson's work, such as honorifics or indirect speech acts, will be considered "polite" only if they go beyond their normal usage as socio-culturally constrained forms of politic behavior. For example, according to Watts, responding to an offer "Would you like some more coffee?" by nonsaliently saying "Yes, please." (Richard J Watts, 2003, p. 186) should be considered a politic behavior. On the other hand, if the response is instead "Yes, please, that's very kind, coffee would be wonderful." Culpeper (2011) is considered to be polite as it is "perceived to go beyond what is expectable" (Richard J Watts, 2003, p. 19). Based on these observations, Locher & Watts (2005) claim that, a better approach than Brown and Levinson's "facework" to encompass all degrees of polite speech will be a larger "relational work" that considers all-ranging levels of politeness from marked im/polite to unmarked politic utterances.

Watts and colleagues' claim can be re-framed as calling for the need to consider wide-ranging informational-social tradeoffs. Watts' distinction between politic versus polite speech lies in whether the addressee recognizes and pays attention to the divergence between speakers' informational and

social goals, as revealed by the discrepancy between literal versus intended meaning, and tries to infer the intended meaning. For politic utterances, the mismatch between literal and intended meaning is unmarked and thus remains unnoticed; For polite utterances, the mismatch is salient and the listener is called to pay attention to the reason for that mismatch, and to infer what a speaker was truly trying to say or what the speaker truly believed (i.e. the true state of the world).

One benefit of the current goal-based account over Watts et al.'s claim is that there is no need to classify different kinds of polite speech as different categories (polite vs. politic). Rather, by observing how subtle changes in the literal meaning, intended meaning, or the true state of the world can lead to changes in a speaker's inferred goals or evaluation of politeness, one can examine the phenomenon as a whole. Watts et al.'s theory makes it rather difficult to know exactly what utterances count as polite vs. politic; for example, should the utterance "This is just a small gift, but I hope you like it" produced in a gift-giving act be considered politic or polite? The goal-based framework obviates this need and quantifies the relative degree of politeness in terms of goal tradeoff.

Other theorists have focused on speakers' attempt to abide by societal obligations and expectations, explicitly distinguishing between their accounts and "strategic" approach to politeness based on individual goals (e.g., P. Brown & Levinson, 1987). Fraser & Nolen (1981) put forth an account for politeness based on the idea of "conversational contract" (CC), which posits that interactants bring a set of rights and obligations to the conversation. Observance of these obligations, according to these authors, are not strategic but rather "getting on with the task in hand in light of the terms and conditions of the CC" (Fraser, 1990, p. 233). Critiques of polite speech analysis based on facework by sociolinguists (Ide, 1989; Mao, 1994; Matsumoto, 1988) assert the importance of considering speakers' desire to fulfill societal obligations. They go against Brown and Levinson's notion of "negative face" as it is heavily loaded with the assumption of individuality as opposed to group identity, which especially becomes prominent in East Asian cultures and languages. Thus, they may argue that an utterance such as "This is just a small gift, but I hope you like it" in East Asian cultures may not reflect an attempt to save the listener's negative face, but to follow convention based on their assigned role (e.g. within a hierarchical social structure).

The proponents of social obligation approach for politeness focus primarily on speakers' prioritization of observance of social norms over informational goals. This approach rejects the assumption of universality of face management principles claimed by Brown and Levinson, and argues that in

some cultures the need for face-work (especially for individuality) may be weak or non-existent. Instead, a speaker may desire to follow and reinforce social norms and expectations, and decide how to balance between that desire and the goal to transfer information. For example, if a professor is to accompany his long-time mentor to a restaurant, when it comes to time to pay, the professor may say “I insist, you should let me pay for this - please treat me next time” despite thinking that he doesn’t actually want to pay the bill, because it is more important to hide his genuine intentions to abide by what is expected of him as a lower-status individual.

The goal-based framework aims to simultaneously acknowledge the importance of these theories identifying more subtle cases of politeness and unify them with the work for face-saving strategies (e.g., P. Brown & Levinson, 1987) they largely argue against. Whereas the social obligation theories attempt to reject the facework-based accounts of polite speech, the current proposal is that speakers consider both social obligation and facework as potentially important social goals, and try to balance between these social goals and epistemic goals to convey their message. Thus, instead of presuming mutual exclusivity of speakers’ goals either to save face or to abide by social rules, the goal-based framework instead assumes that both goals can apply simultaneously, to different degrees depending on the cultural or conversational context.

1.3.4 Politeness in the game-theoretic approach

Finally, there have been recent attempts to analyze polite speech from a game-theoretic perspective, which assumes that individuals interact with each other in an effort to achieve their own goals. A few proponents of game-theoretic approach have analyzed indirect speech as a whole, viewing polite indirect utterances as a subset. Franke & JÄdger (2016) argue that indirect speech used for demands results in greater likelihood of compliance from the other party, as indirect speech suggests higher stakes for the listener in case the speaker’s wants are not fulfilled. For example, a mobster who is trying to take protection money from a restaurant owner could pose a veiled threat by saying: “Your little daughter is very sweet. She goes to the school in Willow Road, I believe” (Franke & JÄdger, 2016, p. 19). This indirect speech is used to communicate that the restaurant owner’s stakes are high because his daughter can get hurt if he does not pay the money, whereas mobster’s stakes are low because he is free to do whatever he wants in the neighborhood. Franke & JÄdger (2016)’s account thus focuses on the reasons to be indirect to send the informational message more

effectively in the interest of the listener’s observance of the speaker’s wants, though the scope of their explanation does not explicitly cover polite speech.

Addressing a broader set of indirect utterances, Pinker, Nowak, & Lee (2008) describe three possible reasons for speaking indirectly: allowing for plausible deniability and thus preventing (legal) responsibility for the intended meaning (“Ocer, is there some way we could take care of the ticket here?”), avoiding emotional costs of mismatch (such as awkwardness) in perceived relationship between speaker and listener, and generating common reference point that is qualitatively different from direct literal meaning even when the intended meaning is clear to the interactants (“Would you like to come up and see my etchings?”).

Other game-theoretic accounts addressed polite indirect speech specifically, highlighting its rationality despite its potential cost. Van Rooy (2003) argues that polite utterances are costly “handicaps,” which incur a social debt and reducing the social status of the speaker, but which are rationally used to communicate good intentions to the addressee. Quinley (2011) observes that politeness is a form of “trust game,” where making a polite request is rational when assumptions of reputation or observation are in place over multiple iterations of conversational turns.

Game-theoretic approach is similar to the current goal-based approach in that the theory is very much built on the notion of tradeoff of goals, but differs in that it is ultimately focused on the speaker’s goals for the self. In game theory, a speaker’s informational goals are centered around informativity, or effective transmission of the message, to ultimately gain the desired compliance from the other party (Franke & JÄdger, 2016). Similarly, a game-theoretic speaker’s social goals involve face management, but less of the intent to save the listener’s face and more to save the speaker’s own face, to avoid being held responsible for disobeying the law or disrespecting someone (plausible deniability) and to prevent high emotional costs misunderstanding common knowledge assumed between speaker and listener (Pinker et al., 2008). The current proposal differentiates between speakers’ interest oriented toward others (i.e., a genuine desire to inform others or save others’ face), and interest towards the speakers themselves (e.g., a desire to *appear* to be a particular way). I formalize the self-interest based on a recursive reasoning about the other-oriented goals: Speakers reason about listeners who imagine speakers to be genuinely informative or kind, and then the speakers produce utterances that can portray themselves to be such helpful individuals.

1.4 Part 3: Empirical work on polite speech in adults

In previous sections, I have recapitulated previous theoretical perspectives to make a unifying proposal that speakers consider tradeoffs between informational and social goals (e.g., prosocial, self-presentational). In this section, I review empirical evidence from adult interpretation, evaluation and production of polite speech that show their consideration of speaker’s informational and social goal tradeoff.

1.4.1 Interpretation of polite speech

Empirical findings show that adults’ interpretation of polite utterance meanings is based on the epistemic-social goal tradeoff considered by the speaker. For example, people’s interpretation of seemingly irrelevant, ambiguous utterances reveal that listeners pay attention to the epistemic-social tradeoff, attributing politeness as a strong reason for apparent irrelevance. T. Holtgraves (1998) examined interpretations of replies to someone asking for an opinion on some performance that he just gave (e.g. “What did you think of my presentation?”), where the replies seemingly did not directly address the questions and involved relevance violations (“It’s hard to give a good presentation”). The author found that people often assign face-threatening meaning to the indirect reply. Additionally, people spent longer coming up with an interpretation when the true state was positive (the addressee actually gave a good presentation) or when the literal meaning of the reply was in the positive direction (“It’s easy to give a good presentation”), which thus made face management an unlikely goal. These findings overall indicate that people identify face management as a key motivation for violating the informativity goal, and struggled when face management was no longer a valid reason for violation of the goal to be informative to the listener.

Another prominent example of polite speech interpretation in line with the idea of the speaker’s goal tradeoff consideration comes from empirical work on cancellation of pragmatic inferences that are normally in force in non-face-threatening, low-stakes contexts. People’s inferences change for non-literal meanings of various quantifiers depending on the presence or severity of face threat toward the listener. For example, the term “possibly” could be used to convey a probability greater than 0 and up to 1 in its literal sense, but usually people assume that it denotes a probability that is neither too high nor too low. This is because people assume that a speaker would use quantifiers in the most efficient, informative way possible (Grice, 1975), and if the speaker wanted to convey a

greater certainty then she would have used a stronger term, such as “certainly.”

However, when the situation involves a potential face threat or a high stake for the listener, the speaker would have a reason to be more vague, to accommodate for face-saving or plausible deniability. Indeed, Bonnefon and colleagues found that people interpret the word “possibly” as implying a greater likelihood when it is used to describe a condition of a higher stake (e.g. “You will possibly suffer from deafness soon” or “Your pain is possibly going to increase”) than a lower stake (“You will possibly suffer from insomnia soon” or “Your pain is possibly going to decrease”; Bonnefon & Villejoubert (2006); Pighin & Bonnefon (2011)).

Similarly, people’s interpretation of other quantifier phrases such as “some” and “A or B” differ depending on the contexts. People usually endow upper-bounded meanings to “some” and “A or B” in non-face-threatening contexts assuming speakers’ epistemic cooperativeness (Breheny, Katsos, & Williams, 2006; e.g., Huang & Snedeker, 2009): “I ate some of the cookies” is interpreted to mean “... some of the cookies but not all,” and “she ate the cake or the salad” to mean “... the cake or the salad but not both.” In situations involving face threat, however, people make different inferences: “some” and “or” in “some of the audiences hated your talk” and “We will cut your salary or take away your company car” are interpreted in their broader sense (“some and possibly all” and “A or B and possibly both”; Bonnefon et al., 2009; Feeney & Bonnefon, 2013). Furthermore, a long pause before the utterance, signaling expectation of bad news to the listener, heightened the effect for the polite interpretation of “some” (Bonnefon, Dahl, & Holtgraves, 2015). In sum, people infer different non-literal intended meanings depending on the context, since the speaker is expected to focus more towards social, face-saving goal than towards epistemic goal in the presence of potential face threats.

1.4.2 Evaluation of polite speakers’ intentions

People’s evaluation of strategies to be polite is also based on the balance of epistemic-social tradeoffs. Clark & Schunk (1980) showed that people rate politeness of requests differently depending on how much the literal meaning of indirect requests would benefit the listener or reduce the cost to the listener. For example, requests that implied that the speaker asked for the listener’s permission (“May I ask you?”) were rated as relatively more polite, whereas those requests that assumed the listener’s obligation to reply to the speaker (and therefore incurs a cost for the listener; “Shouldn’t

you tell me?”) were rated as more impolite. Thus, with the increase of the maintenance of the listener’s face, to be free of imposition or obligation in this case, relative to the degree of message transfer (i.e. access to the intended meaning that was assumed to be constant; e.g. “Tell me the way to Jordan Hall”), people’s judgment for politeness of requests also increased.

Interestingly, hints that seem to prioritize face-saving due to its relatively high degree of indirectness are not always evaluated as most polite. According to P. Brown & Levinson (1987)’s rank ordering of polite strategies, hints (i.e. indirect speech with meanings that are “off the record”) are supposedly more polite compared to other politeness strategies. However, empirical evidence is divided on this issue (e.g., Pinker et al., 2008; Terkourafi, 2002; Yu, 2011). For example, Blum-Kulka (1987) found that people rate conventional indirect requests (e.g. “Would you mind moving your car?”) as more polite than direct orders (“Move your car”) or hints (“We don’t want any crowding”). Along similar veins, among possible replies to a question asking for an opinion on a newly bought dress, people perceive evasive replies (e.g. “It seems like clothes are getting terribly expensive”) to be better than direct utterances (“I don’t think it looks very good on you”) or hints that sound irrelevant (“I’m going to take my vacation next month”; T. Holtgraves, 1986).

One potential reason why hints and irrelevant remarks are not evaluated as best politeness strategies is that they focus too much on face-saving but not information transfer, leading to a poor balance in epistemic-social tradeoff. Indeed, based on the results of her study, Blum-Kulka concluded that politeness is “interactional balance achieved between two needs: The need for pragmatic clarity and the need to avoid coerciveness.” This is in agreement with the current argument that people consider both the epistemic and social goals as important in communication, and the balance between these two goals determines what is optimally “polite”.

1.4.3 Production of polite speech

Adult production of polite utterances reveal their attempt to balance between epistemic and social goals. Adults spontaneously produce requests in polite forms that do not convey their message in the most direct manner (Clark & Schunk, 1980). Even in situations of conflict, people try to balance between their desire to convey their message and desire to control their politeness level; indeed, speakers exhibit politeness strategies even while arguing, preventing unnecessary offense to their interactants (T. Holtgraves, 1997).

Sometimes people decide to compromise the level of unambiguity or certainty that is communicated to the listener due to social reasons. T. Holtgraves & Perdeu (2016) tested what degree of (un)certainty with which people would communicate potentially face-threatening information. They found that if the event referred to was more severe (and therefore there is greater face threat involved) or if the person whose face was under threat was the listener (as opposed a third party) participants used terms with greater uncertainty. Thus, when there is a greater risk for face threat, and therefore more reason to prioritize face management, speakers sacrifice information transfer and convey information more ambiguously.

Situational factors can also lead speakers to prioritize epistemic versus social goals differently. P. Brown & Levinson (1987) claimed that the degree of the face-threat of an act is determined by three factors: the listener's power status with respect to the speaker; the degree of social distance between the speaker and the listener; and the degree of imposition of the act (which may be culturally determined). Thus, a speaker's attempt to minimize face threat would increase as these three factors increase in magnitude. If their claim is true, within the goal-based framework, speakers would prioritize face-saving goal relatively more as the effect of these sociological factors increase.

Indeed, empirical evidence suggest that the lower power status of a speaker, and requests with greater imposition, lead to greater degree of face-saving (Blum-Kulka, Danet, & Gherson, 1985; T. Holtgraves & Yang, 1992; Leichty & Applegate, 1991; Lim & Bowers, 1991). For example, in Lim & Bowers (1991)'s study, participants were asked to write down their probable utterances in an imaginary situation involving a potential face threat and varying degrees of power status of the listener. The results showed that participants produced more "tactful" utterances with more indirect expression of amount of imposition on the listener (e.g. "I'd greatly appreciate it if you could write the paper on behalf of the whole group") in a scenario involving equal power between the speaker and listener than when the speaker had higher power than the listener (e.g. "I know it's imposing a lot, but you gotta write the paper again").

Empirical work on the effect of distance on politeness has been more inconsistent where some report results that are consistent with Brown and Levinson's claim that greater distance leads to higher levels of politeness (T. Holtgraves & Yang, 1992), while others report the opposite (Baxter, 1984; R. Brown & Gilman, 1989). One potential reason for the inconsistency can be that social distance may be confounded with affect (i.e. how much the speaker likes the listener). Nonetheless,

regardless of the directionality, all three variables of power, distance, and degree of imposition affect speakers' politeness level; in other words, speakers consider these factors to determine how much priority they should place on social goal to save the listener's face as opposed to epistemic goal to effectively convey their message.

In sum, empirical work with adults reveal that people interpret, evaluate, and produce polite speech in ways that are consistent with the notion of the epistemic-social goal tradeoff.

1.5 Part 4: Empirical work on polite speech in children

In this section, I argue that children's acquisition of polite speech shows a pattern consistent with the goal-based account: Children show early competence with polite speech that involves minimal conflict between speaker goals, and gradually learn to produce and understand polite speech with more complex patterns of competition between speaker goals.

1.5.1 Rule-based polite speech

Theories on children's sensitivity to conventional rules predict children's early understanding of simple, rule-based politeness (i.e. that does not involve explicit conflict between a speaker's epistemic and social goals). In the domain of morality, Turiel (1977) argued that children's early-emerging understanding of social conventions is primarily based on rules and actions enforced by social authorities, and then the development of a more flexible person-orientation occurs throughout childhood. In other words, children may first start by following simple moral rules such as "do not cheat" or "do not steal," but as they grow older they may become more flexible as they consider more complex circumstances depending on the needs of the individuals and society involved. In the context of polite speech, children may show similar developmental trajectory: they could first faithfully follow rules such as "say please" or "do not lie," but as they become more sensitive to intricate demands of different individuals and contexts, they may start to try to balance between different goals to satisfy those demands. According to this hypothesis, children should start to show mastery of rule-based polite utterances, such as syntactic politeness markers (e.g. "please") and conventional request forms (e.g. "Can you . . ."), earlier than other more complex forms of polite speech (e.g. white lies).

Indeed, children start producing rule-based polite utterances early. From very early on, parents teach children to be polite by following normative rituals to say "please", "thank you", "hello" and

“good-bye” (Gleason, Perlmann, & Greif, 1984). Children start producing the simple polite marker “please” early at 2.5 years (Read & Cherry, 1978) and request forms increase in their variety and frequency with age (E. Bates, 1976; E. Bates & Silvern, 1977; Bock & Hornsby, 1981; Ervin-Tripp, 1982; Nippold, Leonard, & Anastopoulos, 1982).

Young children learn to produce these utterances with appropriate levels of politeness depending on context (Ryckebusch & Marcos, 2004; e.g., Snow, Perlmann, Gleason, & Hooshyar, 1990). For example, children use different forms of requests depending on how they are instructed to make the requests. At three years, children start using different forms of request when they are instructed to “tell” versus “ask” an addressee to give them a puzzle piece (Bock & Hornsby, 1981). Similarly, five-year-olds are able to modify requests depending on whether they are asked to make a request in a nice versus bossy way (Becker, 1986). E. Bates & Silvern (1977) showed that even two-year-olds were able to modify their requests to make them more polite, when they were instructed to make a request to an old lady and then ask her again “in the nicest way possible.” However, children as old as seven fail to differentiate the meaning of “ask” and “tell” when asked Ervin-Tripp (1977) which suggests that production of context-relevant polite speech seems to precede the precise, conscious understanding of goals for speaking politely or impolitely.

Children are also able to adjust the level of politeness based on the listener age and status. Even at two years, children use a polite form of request (e.g. “Can I have...”) to an adult but an imperative form (e.g. “Give me...”) to a peer (Corsaro, 1979; Shatz & Gelman, 1973). Children around 2.5 years were found to use less polite language with their fathers compared to their mothers (Ervin-Tripp, Guo, & Lampert, 1990; Ervin-Tripp, O’Connor, & Rosenberg, 1984). These production behaviors are consistent with Brown and Levinson’s prediction that a speaker’s need for politeness strategy will be heightened as the addressee’s power status increases.

Furthermore, children make polite requests based on the degree of imposition of their demands, and respond sensitively to the resistance to get what they want. By five years, children produce polite speech that matches the cost of the request they make, using more polite forms of requests for apparently more costly impositions (Ervin-Tripp et al., 1990). James (1978) showed that four-year-olds adjusted commands to make them more polite toward older addressees, whereas when they had to make requests for getting what they want (which places higher cost on the speaker if the request is not met) children made the utterances maximally polite regardless of the listener status.

Axia & Baroni (1985) showed that seven-year-old Italian children were able to adjust the polite level of request depending on resistance. Additionally, children in preschool years and older have been observed to repeatedly use the word “please” in their pleading to resistant mothers and peers, as strategies to get what they want (Finley & Humphreys, 1974; Kyratzis & Guo, 2001; Wilson & Wood, 2004). Thus, children use politeness as a strategy to gain compliance for their goal in the intended message.

Comprehension of rule-based polite speech also seems to emerge early, though evidence is more gradual and controversial compared to production. Initial evidence seemed to suggest that producing a request with “please” is judged to be polite by three years of age (E. Bates, 1976; E. Bates & Silvern, 1977). However, in a study by Nippold et al. (1982), this judgment of “please” as polite was only replicated starting at five years of age, but not younger. Furthermore, in the same study, children did not differentiate between requests that adults judged to differ in politeness level (e.g. “Give me some candy, please” vs. “Can you give me some candy, please?”) until seven years of age. But these initial studies have a few unresolved issues, including the lack of statistical tests to assess each age group’s performance and lack of systematic manipulation of cues other than syntax (e.g., voicing or facial expressions). Thus, more evidence is required to confirm early understanding of conventional polite requests, but preliminary evidence so far suggests that children start to show some understanding of these requests before preschool years.

1.5.2 Indirect speech

Now I turn to children’s production and understanding of polite speech that is more complex than simple rule-based utterances and reflects a strategic tradeoff between speaker goals. First, what do we know about children’s understanding of indirect speech, where informativity is compromised for social goals?

Despite the wealth of studies looking at conventional indirect requests and simple polite markers (“please”), very few studies have looked at children’s understanding of non-conventional indirect requests (e.g. hints); there is only limited evidence that comprehension of under-informative indirect requests is more difficult and is acquired at a later age than conventional requests. Bernicot and colleagues have tested French-speaking children’s abilities to evaluate implications of requests for which the intended meaning is distinct from the literal meaning. Bernicot & Legros (1987) tested children’s

ability to understand implications of directives in different forms: direct directives (e.g. “Give me the spade”) versus nonconventional indirect directives (“I can’t make a castle with my hands”). They found that even by six years, children have difficulty judging what would be an appropriate reaction by the speaker when the addressee does not comply with the directives. Bernicot, Laval, & Chaminaud (2007) found some evidence that by eight years, children are able to infer appropriate responses to nonconventional requests in the form of hints. Finally, ten-year-old children, similar to adults (as shown in Blum-Kulka, 1987), judge that conventional indirect requests are the “right thing to say” more often compared to direct orders or hints, whereas five- and seven-year-olds do not show this metapragmatic reasoning (Bernicot, 1991).

1.5.3 White lies

Next I consider another category of polite utterances reflecting speakers’ goal tradeoff: white lies. Children start producing white lies early on. Talwar and colleagues have shown the earliest evidence of white lie-telling in children: They found that the majority of three- to seven-year-old child participants in their studies told white lies and stated that an adult “looks okay for the picture” even though she has a conspicuous mark of lipstick on her nose (Talwar & Lee, 2002), or lied to a gift-giver about her gift that they actually found undesirable (Talwar, Murphy, & Lee, 2007).

Though they produce white lies early, young children’s understanding of goals for production of white lies seems limited. When the participants in Talwar and colleagues’ studies were asked for reasons why they decided to tell white lies, many stated they did not know why. Similarly, Italian children of same age group struggled to predict what the protagonist of a story might say in a situation where telling truth would not be polite (Airenti & Angeleri, 2011). Although there was an increase in the proportion of correct responses with age, even six-year-olds did not predict the correct response above chance unless they were given an additional cue to think about politeness-related reasons. Thus, despite the evidence of early spontaneous production of white lies in familiar situations, children seem to have limited understanding of goals behind white lie-telling.

Studies of comprehension and evaluation of white lies show that young children are more charitable toward prosocial lies compared to other kinds of lies. Evidence suggests that by four years, children perceive prosocial lies as better than malicious lies (Bussey, 1999), and they rate lies as nicer than truths in politeness situations (M. Song & Song, 2014). G. D. Heyman, Sweet, & Lee

(2009) observed that 7- to 11-year-old children rated lie-telling more favorably in politeness situations (e.g. a teacher gave the protagonist an undesirable gift) than in transgression situations (e.g. the protagonist damaged a library book). This tendency to rate prosocial lies positively increases with age (Walper & Valtin, 1992).

Children's favorable perception of white lies tends to depend on their recognition of speaker goals and reasons for white lie-telling. When G. D. Heyman et al. (2009) asked participants for explanations of their evaluation of the protagonist's behavior (telling a lie or truth) in politeness context, the participants were more likely to rate lie-telling more favorably and truth-telling more negatively when their focus was on the impact of the statement on others rather than the veracity of statements. Studies in other countries have also found that children in elementary school years view lie-tellers more positively if the lie is told in a public, face-threatening situation (Ma, Xu, Heyman, & Lee, 2011), or is intended to benefit others rather than the speaker herself (Fu, Heyman, Chen, Liu, & Lee, 2015). Overall these findings show children evaluate white lies and blunt truths differently based on which communicative goal or intention is prominent in the given context.

As they get older, children make more subtle inferential judgments about implications of white lies, which shows their ability to identify epistemic and social goals as separate intentions manifested in white lies. Children become able to differentiate between various traits of a speaker based on her goal tradeoff decision. Xu et al. (2013) told 7- 9- and 11-year-old Chinese children stories of lie-tellers and truth-tellers with either intent of helping or harming the addressee, and looked at their ratings of benevolence and trustworthiness of the speaker. The participants rated helpful characters (both lie-tellers and truth-tellers) as nice and harmful characters as mean; for trust evaluations, younger children relied more on honesty whereas older children relied more on the intentions. This finding reveals an interesting trend in which children of younger ages are more narrowly focused on the speaker's truth-telling goal, whereas older children consider the goal tradeoff more holistically and gradually shift their preference toward the face-saving goal. Indeed, in another study by Xu, Bao, Fu, Talwar, & Lee (2010), Chinese children from the same age group who considered both honesty and politeness issues in evaluating white lie-tellers tended to tell a white lie more often than those who focused on only one of the two issues.

1.5.4 Summary

In sum, the development of production and understanding of polite speech show predicted trajectories based on the currently proposed goal-based account. A greater degree of conflict or tradeoff between the speaker's informational and social goals leads to more challenging cases of polite speech that are more difficult to produce or comprehend. Thus, from early on, children produce and understand many instances of rule-based polite speech using simple markers of politeness, involving minimal conflict between epistemic and social goals. As for indirect speech and white lies, cases that reflect speakers' decisions based on the goal tradeoff, younger children before elementary school years show only limited understanding of motivation behind these more complex utterances. Finally, school-year children become competent at identifying reasons for both truth-telling (for epistemic reasons) and prosocial lie-telling (for social reasons) and making informed evaluations of speakers based on their tradeoff decisions.

1.6 Conclusions

In this Chapter I argued for a goal-based account of polite speech, in which speakers consider their competing social goals (informational, pro-social, and self-presentational goals) to speak politely. The goal-based approach provides a way to unify previous theoretical frameworks for polite speech, and is consistent with empirical evidence in adults' and children's production and comprehension of polite speech.

In Chapters 2 through 4, I will present empirical evidence from adults and children to support the goal-based account of polite speech. Chapter 2 presents a formal model of how polite speech emerges. We explored how speakers are expected to speak politely in situations where they are asked to provide feedback for the addressee's performance (e.g., poem recital). We built a computational model based on the assumption that speakers should consider goals to be informational, pro-social and self-presentational in speaking politely (as addressed in Chapter 1), and we show that our model successfully captures important patterns of adult predictions for polite speech (e.g., white lies and indirect speech).

Chapters 3 and 4 examine children's understanding of polite speech. In Chapter 3, we look at whether 2- to 4-year-old children are able to understand that speakers account for prosocial goals in

their speech, by examining their evaluation of polite requests (e.g., “Can you please pour me more water?”). In Chapter 4, we investigate whether children reason about tradeoff between different goals (e.g., informativity vs. prosociality) by looking at their evaluation of white lies (e.g., “Your cookie was tasty”) versus blunt truths (“Your cookie was yucky”).

The primary goal of this dissertation, then, is to probe whether and how adults and children reason about polite speech as reflecting competing social goals, such as goals to be informative, to be kind to others, and to present oneself in a good light, and tradeoffs between these goals.

Chapter 2

Adults consider tradeoffs between competing social goals to predict polite language use¹

Language is a remarkably efficient tool for transmitting information. Yet human speakers make statements that are inefficient, imprecise, or even contrary to their own beliefs, all in the service of being polite. What rational machinery underlies polite language use? In this Chapter, I present evidence that polite speech emerges from the competition of three communicative goals: to convey information, to be kind, and to present oneself in a good light. We formalized this goal tradeoff using a probabilistic model of utterance production, which predicts human utterance choices in socially-sensitive situations with high quantitative accuracy, and I show that our full model is superior to its variants with subsets of the three goals.

2.1 Introduction

We rarely say exactly what’s on our mind. Although “close the window!” could be an effective message, we dawdle by adding “can you please...?” or “would you mind...?” Rather than tell an

¹This chapter is submitted and currently under review at *Open Mind*, and is joint work with Michael Henry Tessler, Noah D. Goodman and Michael C. Frank.

uncomfortable truth, socially-aware speakers lie (“Your dress looks great!”) and prevaricate (“Your poem was so appropriate to the occasion”). Such language use is puzzling for classical views of language as information transfer (B  ijhler, 1934; Frank & Goodman, 2012; Jakobson, 1960; Shannon, 1948). On the classical view, transfer ought to be efficient and accurate: Speakers are expected to choose succinct utterances to convey their beliefs (Grice, 1975; Searle, 1975), and the information conveyed is ideally truthful to the extent of a speaker’s knowledge. Polite speech violates these basic expectations about the nature of communication: It is typically inefficient and underinformative, and sometimes even outright false. Yet even young speakers spontaneously produce requests in polite forms (Axia & Baroni, 1985), and adults use politeness strategies while arguing (T. Holtgraves, 1997), even though polite utterances may risk high-stakes misunderstandings (Bonnefon, Feeney, & De Neys, 2011).

If politeness only gets in the way of effective information transfer, why be polite? Clearly, there are social concerns, and most linguistic theories assume utterance choices are motivated by these concerns, couched as either polite maxims (Leech, 1983), social norms (Ide, 1989), or aspects of a speaker and/or listener’s identity, known as *face* (P. Brown & Levinson, 1987; Goffman, 1967). Face-based theories predict that when a speaker’s intended meaning contains a threat to the listener’s face or self-image (and potentially the speaker’s face), her messages will be less direct, less efficient, and possibly untruthful. Indeed, listeners readily assume speakers’ intentions to be polite when interpreting utterances in face-threatening situations (Bonnefon et al., 2009). How this socially-aware calculation unfolds, however, is not well understood. When should a speaker decide to say something false (“Your poem was great!” based on an example from Bonnefon et al. (2009)) rather than just be indirect (*Some of the metaphors were tricky to understand.*)? How does a speaker’s own self-image enter into the calculation?

We propose a utility-theoretic solution to the problem of polite language use by quantifying the tradeoff between competing communicative goals. In our model, speakers attempt to maximize utilities that represent their communicative goals: informational utility—derived via classical, effective information transmission; social utility—derived by being kind and saving the listener’s face; and self-presentational utility—the most novel component of our model, derived by appearing in a particular way to save the speaker’s own face. Speakers then produce an utterance on the basis of its expected utility (including their cost to speak). The lie that a poem was great provides social

utility by making the writer feel good, but does not provide information about the true state of the world. Further, if the writer suspects that the poem was in fact terrible, the speaker runs the risk of being seen as uncooperative.

We assume that speakers’ utilities are weighed within a probabilistic model of pragmatic reasoning: the Rational Speech Act (RSA) framework (Frank & Goodman, 2012; N. D. Goodman & Frank, 2016a). Speakers are modeled as agents who choose utterances by reasoning about their potential effects on a listener, while listeners infer the meaning of an utterance by reasoning about speakers and what goals could have led them to produce their utterances. This class of models has been effective in understanding a wide variety of complex linguistic behaviors, including vagueness (Lassiter & Goodman, 2017), hyperbole (Kao, Wu, Bergen, & Goodman, 2014), and irony (Kao & Goodman, 2015), among others. In this framework, language use builds on the idea that human social cognition can be approximated via reasoning about others as rational agents who act to maximize their subjective utility (Baker, Saxe, & Tenenbaum, 2009), a hypothesis which has found support in a wide variety of work with both adults and children (e.g., Jara-Ettinger, Gweon, Schulz, & Tenenbaum, 2016; S. Liu, Ullman, Tenenbaum, & Spelke, 2017). RSA models are defined recursively such that speakers S reason about listeners L , and vice versa. We use a standard convention in indexing and say a pragmatic listener L_1 reasons about what intended meaning and goals would have led a speaker S_1 to produce a particular utterance. Then S_1 reasons about a *literal listener* L_0 , who is modeled as attending only to the literal meanings of words (rather than their pragmatic implications), and hence grounds the recursion. The target of our current work is a model of a polite speaker S_2 who reasons about what to say to L_1 by considering informational, social, and self-presentational goals (Figure 2.1).

We evaluate our model’s ability to predict human utterance choices in situations where polite language use is expected. Imagine Bob recited a poem and asked Ann how good it was. Ann (S_2) produces an utterance w based on the true state of the world s (i.e., the rating, in her mind, truly deserved by Bob’s poem) and a set of goal weights $\hat{\phi}$, that determines how much Ann prioritizes each of the three possible goals. Ann’s production decision is softmax, which interpolates between maximizing and probability matching (via λ_{S_2} ; N. D. Goodman & Stuhlmüller, 2013):

$$P_{S_2}(w|s, \hat{\phi}) \propto \exp(\lambda_{S_2} \cdot \mathbb{E}[U_{total}(w; s; \hat{\phi}; \phi_{S_1})]).$$




Figure 2.1: Diagram of the model: The polite speaker observes the true state and determines her goal between three utilities (informational, social, and presentational), and produces an utterance.

We posit that a speaker's utility contains three distinct components: informational, social, and presentational. The total utility U_{total} of an utterance is thus the weighted combination of the three utilities minus the utterance cost $C(w)$:

$$U_{total}(w; s; \hat{\phi}; \phi_{S_1}) = \phi_{inf} \cdot U_{inf}(w; s) + \phi_{soc} \cdot U_{soc}(w) + \phi_{pres} \cdot U_{pres}(w; \phi_{S_1}) - C(w).$$

We define *social utility* (U_{soc}) as the expected subjective utility of the state $V(s)$ implied to the pragmatic listener by the utterance: $U_{soc}(w) = \mathbb{E}_{P_{L_1}(s|w)}[V(s)]$. The subjective utility function $V(s)$ could vary by culture and context; we test our model when states are explicit ratings (e.g., on a 4-point scale) and we assume a positive linear value relationship between states and values V to model a listener’s preference to be in a highly rated state (e.g., Bob would prefer to have written a poem deserving 4 points rather than 1 point).

At the same time, a speaker may desire to be epistemically helpful, modeled as standard *informational utility* (U_{inf}). The informational utility indexes the utterance’s *surprisal*, or amount of information the listener (L_1) would still not know about the state of the world s after hearing the speaker’s utterance w (e.g., how likely is Bob to guess Ann’s actual opinion of the poem): $U_{inf}(w) = \ln(P_{L_1}(s|w))$. Speakers who optimize for informational utility produce accurate and informative utterances while those who optimize for social utility produce utterances that make the listener feel good.

If a listener is uncertain how their particular speaker is weighing the competing goals to be honest vs. kind (informational vs. social utilities), he might try to infer the weighting (e.g., “was she just being nice?”). But a sophisticated speaker can produce utterances in order to appear *as if* she had certain goals in mind, for example making the listener think that the speaker was being both kind and informative (“she wanted me to know the truth but without hurting my feelings”). The extent to which the speaker *appears* to the listener to have a particular goal in mind (e.g., to be kind) is the utterance’s *presentational utility* (U_{pres}). The speaker gains presentational utility when her listener believes she has particular goals, represented by a mixture weighting ϕ_{S_1} between trying to be genuinely informative vs. kind. Formally,

$$U_{pres}(w; \phi_{S_1}) = \ln(P_{L_1}(\phi_{S_1} | w)) = \ln \int_s P_{L_1}(s, \phi_{S_1} | w).$$

The speaker conveys a particular weighting of informational vs. social goals (ϕ_{S_1}) by considering the beliefs of listener L_1 , who hears an utterance and jointly infers the speaker’s utilities and the true state of the world:

$$P_{L_1}(s, \phi_{S_1} | w) \propto P_{S_1}(w | s, \phi_{S_1}) \cdot P(s) \cdot P(\phi_{S_1}).$$

The presentational utility is the highest-order term of the model, defined only for a speaker thinking about a listener who evaluates a speaker (i.e., defined for S_2 , but not S_1). Only the social and informational utilities are defined for the S_1 speaker (via reasoning about L_0); thus, S_1 's utility weightings can be represented by a single number, the mixture parameter ϕ_{S_1} . Definitions for S_1 and L_0 otherwise mirror those of S_2 and L_1 and can be found in the Supplementary Materials: Model details section.

Finally, more complex utterances incur a greater cost, $C(w)$ – capturing the general pressure towards economy in speech. In our work, utterances with negation (e.g., *not terrible*) are assumed to be slightly costlier than their equivalents with no negation (this cost is inferred from data; see Supplementary Materials).

Within our experimental domain, we assume there are four possible states of the world corresponding to the value placed on a particular referent (e.g., the poem the speaker is commenting on), represented in terms of numbers of hearts (Figure 2.1): $S = s_0, \dots, s_3$. Since the rating scale is relatively abstract, we assume a uniform prior distribution over possible states of the world. The set of utterances is $\{\textit{terrible}, \textit{bad}, \textit{good}, \textit{amazing}, \textit{not terrible}, \textit{not bad}, \textit{not good}, \textit{and not amazing}\}$. We implemented this model using the probabilistic programming language WebPPL (N. D. Goodman & Stuhlmüller, 2014) and a demo can be found at <http://forestdb.org/models/politeness.html>.

2.2 Model predictions

The pragmatic listener model L_1 draws complex inferences about both the true state of the world (Fig. 2.2A) and the speaker's goals (Fig. 2.2B). Upon hearing *[Your poem] was terrible* (Fig. 2.2A and 2.2B top-left), the listener infers the poem is probably truly terrible (i.e., worthy of zero hearts) and that the speaker has strong informational goals. *It was amazing* is more ambiguous (Figure 2.2A and 2.2B top-right): The poem could indeed be worthy of three hearts, but it is also plausible the speaker had strong social goals and the poem was mediocre. Negation makes the meanings less precise and introduces more uncertainty into the inference about the state: A listener who hears *It wasn't amazing* sees it as a relatively kind way of saying that the poem was quite bad (0 or 1 hearts), inferring a balance of social and informational goals for the speaker (Figure 2.2A and 2.2B bottom-right). *It wasn't terrible* is the most open-ended, leaving open the possibility that the poem

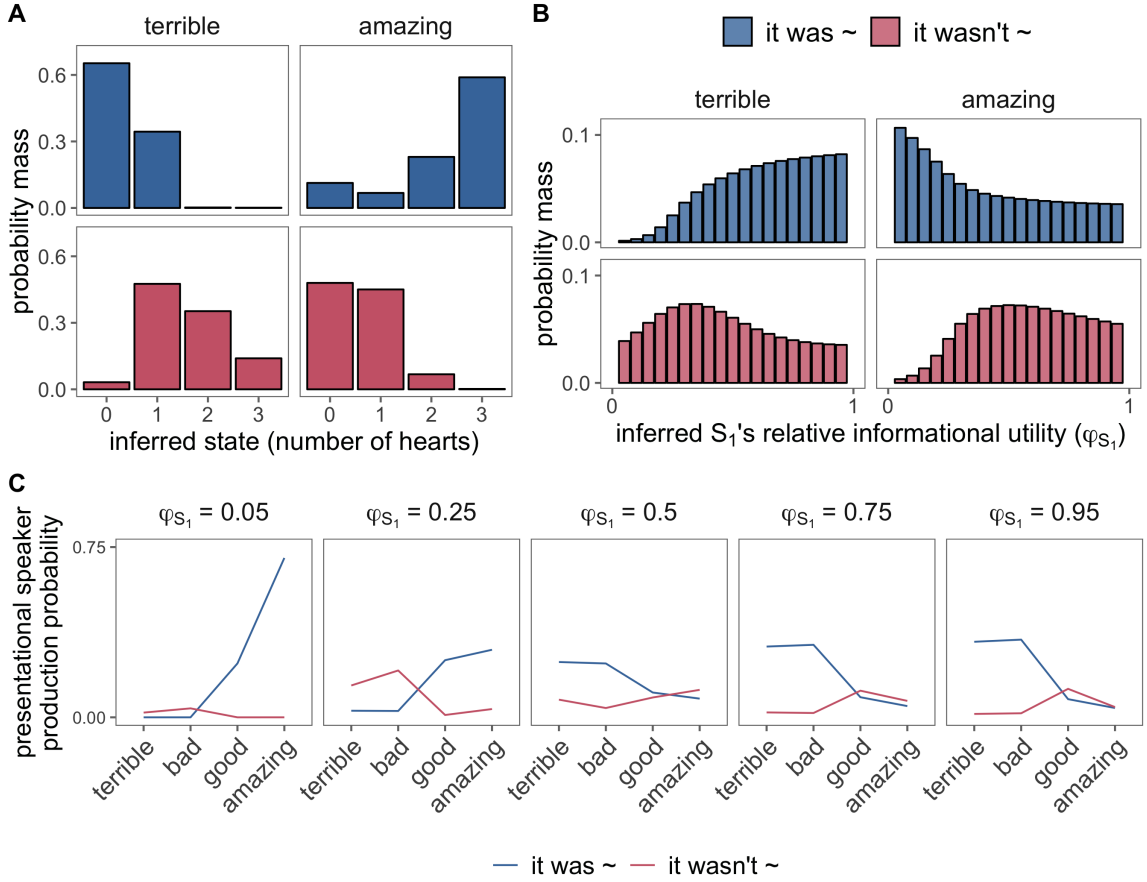


Figure 2.2: Model behavior. Listener inferences about the true state (e.g., the rating truly deserved by the poem; A) and the speaker's utility weighting (ϕ_{S_1} or how informational vs. social the speaker is, where $\phi_{S_1} = 0$ is fully social, and $\phi_{S_1} = 1$ is fully informational; B) as a function of the utterance heard (facets). C: Purely self-presentational speaker production behavior as a function of the kind of speaker they wish to present themselves as (facets; relatively more informational, e.g., $\phi_{S_1} = 0.05$, vs. social as represented, e.g., $\phi_{S_1} = 0.95$).

was worthy of 0 hearts (i.e., *it was terrible*) but conveying to the listener that the speaker cares about both informational and social goals, with a slight preference of towards being social (Figure 2.2A and 2.2B bottom-left).

The self-presentational utility guides the speaker S_2 to care about how she will be viewed in the eyes of the listener L_1 (Figure 2.2C). If the speaker wants to present herself as someone who is socially-minded (e.g., informational mixture or ϕ_{S_1} of 0.05), she should produce direct, positive utterances (e.g., *amazing*). The best way to appear honest (e.g., informational mixture of 0.95) is


to say direct, negative utterances (e.g., *terrible*). The desire to appear as someone concerned with telling the truth while also caring about the listener's feelings (e.g., ϕ_{S_1} of 0.25) leads the speaker to produce indirect utterances (e.g., *not terrible*). Such indirect speech acts are sufficiently open-ended to include the possibility that the poem was good, but the avoidance of a more direct utterance (e.g., *good*) provides the listener with a way to recover the true state (e.g., the poem was mediocre) by way of reasoning that the speaker cares about his feelings by not saying the blunt truth.

2.3 Experiment: Speaker production task

We made a direct, fully pre-registered test of our speaker production model and its performance in comparison to a range of alternative models, by instantiating our running example in an online experiment.

Imagine that Fiona filmed a movie, but she didn't know how good it was. Fiona approached Yvonne, who knows a lot about movies, and asked "How was my movie?"

Here's how Yvonne **actually** felt about Fiona's movie, on a scale of 0 to 3 hearts:



If Yvonne wanted to BOTH make Fiona feel good AND give accurate and informative feedback,

what would Yvonne be most likely to say?

"It "

Figure 2.3: Example of a trial in the speaker production task.

2.3.1 Participants

202 participants with IP addresses in the United States were recruited on Amazon’s Mechanical Turk.

2.3.2 Design and Methods

Participants read scenarios with information on the speaker’s feelings toward some performance or product (e.g., a poem recital; *true state*), on a scale from zero to three hearts (e.g., one out of three hearts). For example, one trial read: *Imagine that Bob gave a poem recital, but he didn’t know how good it was. Bob approached Ann, who knows a lot about poems, and asked “How was my poem?”* Additionally, we manipulated the speaker’s goals across trials: to be *informative* (“give accurate and informative feedback”); to be *kind* (“make the listener feel good”); or to be *both* informative and kind simultaneously. We hypothesized that each of the three experimentally-induced goals would induce a different tradeoff between social and informational utilities in our model, as well as modulating the self-presentational component. In a single trial, each scenario was followed by a question asking for the most likely produced utterance by Ann. Participants selected one of eight possible utterances, by choosing between *It was* vs. *It wasn’t* and then among *terrible*, *bad*, *good*, and *amazing*.

Each participant read twelve scenarios, depicting every possible combination of the three goals and four states. The order of context items was randomized, and there were a maximum of two repeats of each context item per participant. Each scenario was followed by a question that read, “If Ann wanted to make Bob feel good but not necessarily give informative feedback (or to give accurate and informative feedback but not necessarily make Bob feel good, or BOTH make Bob feel good AND give accurate and informative feedback), what would Ann be most likely to say?” Participants indicated their answer by choosing one of the options on the two dropdown menus, side-by-side, one for choosing between *It was* vs. *It wasn’t* and the other for choosing among *terrible*, *bad*, *good*, and *amazing*.

2.3.3 Behavioral results

Our primary behavioral hypothesis was that speakers describing bad states (e.g., poem deserving 0 hearts) with goals to be both informative and kind would produce more indirect, negative utterances (e.g., *It wasn’t terrible*). Such indirect speech acts both save the listener’s face and provide some

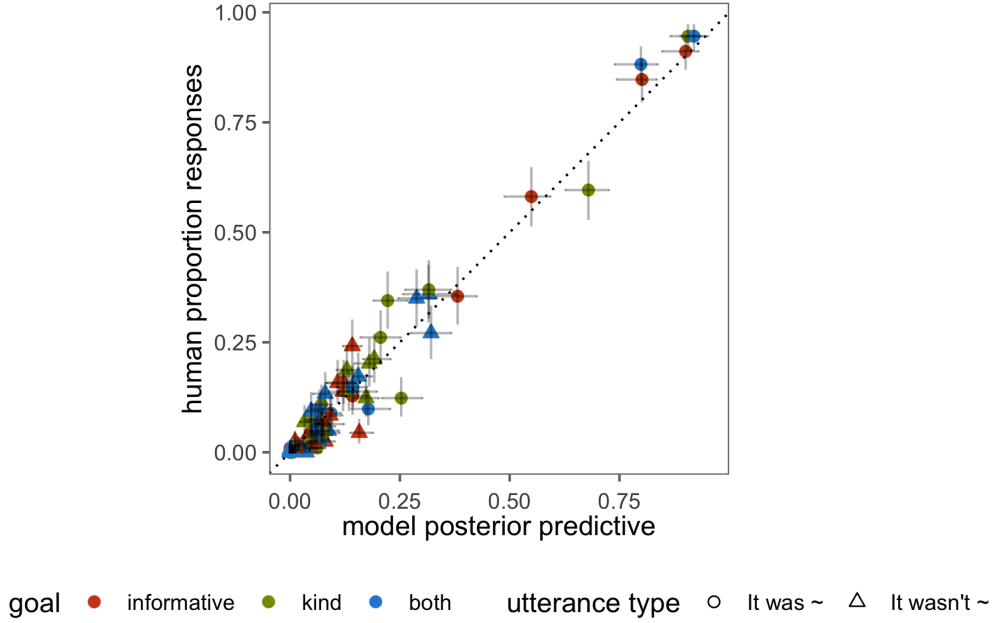


Figure 2.4: Full distribution of human responses vs. model predictions. Error bars represent 95% confidence intervals for the data (vertical) and 95% highest density intervals for the model (horizontal).

information about the true state, and thus, are what a socially-conscious speaker would say (Figure 2.2). This prediction was confirmed, as a Bayesian mixed-effects model predicts more negation as a function of true state and goal via an interaction: A speaker with both goals to be informative and kind produced more negation in worse states compared to a speaker with only the goal to be informative ($M = -1.33$, $[-1.69, -0.98]$) and goal to be kind ($M = -0.5$, $[-0.92, -0.07]$). Rather than eschewing one of their goals to increase utility along a single dimension, participants chose utterances that jointly satisfied their conflicting goals by producing indirect speech.

2.3.4 Model results

The model parameters (softmax parameters and each goal condition's utility weights) can be inferred from the behavioral data using Bayesian data analysis (M. D. Lee & Wagenmakers, 2014). To approximate the literal meanings (i.e., the semantics) of the words as interpreted by the literal listener L_0 , we obtained literal meaning judgments from an independent group of participants (See Supplementary Materials: Literal semantic task section). The posterior predictions from the the




Figure 2.5: Comparison of predictions for proportion of utterances chosen by pragmatic speaker from possible model variants (left) and human data (rightmost) for average proportion of negation produced among all utterances, given true state of 0 heart (on a scale of 0 to 3) and speaker with both goals to be informative and kind. Gray dotted line indicates chance level at 12.5%.

three-utility polite speaker model (informational, social, presentational) showed a very strong fit to participants' actual utterance choices ($r^2(96) = 0.971281$; Figure 2.4). We compared these to six model variants containing subsets of the three utilities in the full model. Both the variance explained and marginal likelihood of the observed data were the highest for the full model (Table 2.1). Only the full model captured participants' preference for negation when the speaker wanted to be informative and kind about truly bad states, as hypothesized (Figure 2.5). In sum, the full set of informational, social, and presentational were required to fully explain participants' utterance choices. The utility weights inferred for the three-utility model (Table 2.2) provide additional in-

Table 2.1: Comparison of variance explained for each model variant and log Bayes Factors quantifying evidence in favor of alternative model in comparison.

model	variance explained	log BF
informational, social, presentational	0.97	—
informational, presentational	0.96	-11.14
informational, social	0.92	-25.06
social, presentational	0.23	-864
presentational only	0.23	-873.83
social only	0.22	-885.52
informational only	0.83	-274.89

Table 2.2: Inferred phi parameters from all model variants with more than one utility.

model (utilities)	goal	ϕ_{inf}	ϕ_{soc}	ϕ_{pres}	ϕ_{S_1}
informational, social, presentational	both	0.36	0.11	0.54	0.36
informational, social, presentational	informative	0.36	0.02	0.62	0.49
informational, social, presentational	social	0.25	0.31	0.44	0.37
informational, presentational	both	0.64	–	0.36	0.17
informational, presentational	informative	0.77	–	0.23	0.33
informational, presentational	social	0.66	–	0.34	0.04
informational, social	both	0.54	0.46	–	–
informational, social	informative	0.82	0.18	–	–
informational, social	social	0.39	0.61	–	–
social, presentational	both	–	0.38	0.62	0.55
social, presentational	informative	–	0.35	0.65	0.75
social, presentational	social	–	0.48	0.52	0.66

sight into how polite language use operates in our experimental context and possibly beyond: *Being kind* (“social”) requires not only weights on social and presentational utilities but equal weights on all three utilities, indicating that informativity is a part of language use even when it is explicitly not the goal. *Being informative* (“informative”) pushes the weight on social utility (ϕ_{soc}) close to zero, but the weight on *appearing kind* (ϕ_{pres}) stays high, suggesting that speakers are expected to manage their own face even when they are not considering others’. *Kind and informative* (“both”) speakers emphasize informativity slightly more than kindness. In all cases, however, the presentational utilities have greatest weight, suggesting that managing the listener’s inferences about oneself was integral to participants’ decisions in the context of our communicative task. Overall then, our condition manipulation altered the balance between these weights, but all utilities played a role in all conditions.

2.4 Discussion

Politeness is puzzling from an information-theoretic perspective. Incorporating social motivations adds a level of explanation, but so far such intuitions and observations have resisted both formalization and precise testing. We present a utility-theoretic model of language use that captures the interplay between competing informational, social, and presentational goals, and provide pre-registered experimental evidence that confirmed its ability to capture human judgments, unlike

comparison models with only a subset of the full utility structure.

To estimate precisely choice behavior in the experiment, it was required to abstract away from natural interactions in a number of ways. Human speakers have access to a potentially infinite set of utterances to select from in order to manage the three-utility tradeoff (*It's hard to write a good poem, That metaphor in the second stanza was so relatable!*). In theory, each utterance will have strengths and weaknesses relative to the speaker's goals, though computation in an unbounded model presents technical challenges (perhaps paralleling the difficulty human speakers feel in finding the right thing to say in a difficult situation; see N. D. Goodman & Frank, 2016a).

For a socially-conscious speaker, managing listeners' inferences is a fundamental task. Our work extends previous models of language beyond standard informational utilities to address social and self-presentational concerns. Further, our model builds upon the theory of politeness as face management (P. Brown & Levinson, 1987) and takes a step towards understanding the complex set of social concerns involved in face management. Our approach can provide insight into a wide range of social behaviors beyond speech by considering utility-driven inferences in a social context (Baker, Jara-Ettinger, Saxe, & Tenenbaum, 2017; Hamlin, Ullman, Tenenbaum, Goodman, & Baker, 2013) where agents need to take into account concerns about both self and others.

Previous game-theoretic analyses of politeness have either required some social cost to an utterance (e.g., by reducing one's social status or incurring social debt to one's conversational partner; Van Rooy, 2003) or a separately-motivated notion of plausible deniability (Pinker et al., 2008). The kind of utterance cost for the first type of account would necessarily involve higher-order reasoning about other agents, and may be able to be defined in terms of the more basic social and self-presentational goals we formalize here. A separate notion of plausible deniability may not be needed to explain most politeness behavior, either. Maintaining plausible deniability is in one's own self-interest (e.g., due to controversial viewpoints or covert deception) and goes against the interests of the addressee; some amount of utility dis-alignment is presumed by these accounts. Politeness behavior appears present even in the absence of obvious conflict, however: In fact, you might be even more motivated to be polite to someone whose utilities are more aligned with yours (e.g., a friend). In our work here, we show that such behaviors can in fact arise from purely cooperative goals (P. Brown & Levinson, 1987), though in cases of genuine conflict, plausible deniability likely plays a more central role in communication.

Utility weights and value functions in our model could provide a framework for a quantitative understanding of systematic cross-cultural differences in what counts as polite. Cross-cultural differences in politeness could be a product of different weightings within the same utility structure. Alternatively, culture could affect the value function V that maps states of the world onto subjective values for the listener (e.g., the mapping from states to utilities may be nonlinear and involve reasoning about the future). Our formal modeling approach with systematic behavior measurements provides an avenue towards understanding the vast range of politeness practices found across languages.

Politeness is only one of the ways language use deviates from purely informational transmission. We flirt, insult, boast, and empathize by balancing informative transmissions with goals to affect others' feelings or present particular views of ourselves. Our work shows how social and self-presentational motives are integrated with informational concerns more generally, opening up the possibility for a broader theory of social language. In addition, a formal account of politeness moves us closer to courteous computation – to machines that can talk with tact.

Chapter 3

Children understand social goals behind polite requests¹

In the last chapter, I examined adults’ understanding of polite speech based on goal tradeoffs. What do children understand about polite speech? Looking at children’s polite speech comprehension can help examine children’s pragmatic understanding more generally, and can be informative for caregivers who want to teach children what it means to be polite. Even though children start to produce polite speech from early on, there is little known about whether they understand intentions behind polite language. In this Chapter, I show that by 3 years, English-speaking preschool children understand that it is more polite and nicer (and less rude and mean) to use politeness markers such as “please” when making requests, and by 4 years, they understand that the use of these politeness markers indicates that the speaker is more socially likeable and is more likely to gain compliance from their conversational partners. This work can help lay the foundation for future work on children’s understanding of polite speech and pragmatic development more generally.

3.1 Introduction

We use and hear polite speech on a daily basis: polite utterances range from simple words of apology (“sorry”) or gratitude (“thanks”) to compliments (“I love your dress!”) and requests (“Can you

¹This chapter is submitted and currently under review for the 41st Annual Meeting of the Cognitive Science Society, and is joint work with Michael C. Frank.

please open the window?”). Yet polite utterances are seemingly inefficient and even misinformative: speakers say “Can you please . . .” when it should suffice to say, “Open the window.” These facts are a mystery for frameworks which describe communication in terms of efficient information transfer (e.g., B  ijhler, 1934; N. D. Goodman & Stuhlm  ijller, 2013; Shannon, 1948): If language is a tool for transferring information, speakers should be as efficient as possible in their communication to prioritize informativity. Nonetheless, everyday politeness is ubiquitous in everyday language use, and adults tend to use strategies to be polite even while arguing (T. Holtgraves, 1997).

So why do people speak politely? Linguistic theories assume that people’s utterance choices are motivated by social concerns, framed as either maxims (Leech, 1983), social norms (Ide, 1989), or listener’s and/or speaker’s public identity (*face*; P. Brown & Levinson, 1987). For example, P. Brown & Levinson (1987)’s theory predicts that if a speaker’s intended meaning contains a threat to the listener’s face or self-image, the speaker’s utterance will be less direct and less informative. For example, if a speaker considered that saying “Open the window” will give the impression that she is in a position to give orders to the listener, she could instead say “Can you please open the window?”, using a more indirect form of request to give the other person a sense of autonomy or freedom from imposition (Clark & Schunk, 1980). Thus, while it may hinder the goal of efficient information transfer, using polite speech can help the speaker save the listener’s face while simultaneously communicating her own positive social goals (Erica J. Yoon, Tessler, Goodman, & Frank, 2017).

Do children speak politely, and if so, what do they understand about polite speech? Previous research shows that children begin producing polite speech early on; They produce “please” at 2.5 years (Read & Cherry, 1978), and request forms increase in their variety and frequency with age (E. Bates, 1976; E. Bates & Silvern, 1977; Bock & Hornsby, 1981; Ervin-Tripp, 1982; Nippold et al., 1982). Young children learn to produce different forms of requests depending on context: For example, by three years children are able to vary their utterances based on whether they are instructed to “tell” versus “ask” an addressee to give them a puzzle piece (Bock & Hornsby, 1981). And even at two years, children are able to modify their requests to make them more polite (“ask in the nicest way possible”; E. Bates & Silvern, 1977). Hence, children’s production of polite speech seems to parallel adult speakers’ desires to produce utterances with appropriate levels of face-saving.

While children appear to produce polite speech from an early age, less is known about whether

they *understand* polite speech. Examining children's comprehension of polite speech is important for a number of reasons. First, children's polite speech understanding can reveal their inferential abilities underlying more general pragmatic understanding: going beyond what was literally said to infer what was intended. For example, children need to understand that, in saying "can you open the window?" the speaker does not literally question the listener's ability to open the window but rather wants to make a polite request. Thus, understanding what children comprehend about polite speech can help see how children are able to infer speaker's intentions behind utterances.

Second, understanding polite speech can have practical implications for education, as caregivers often care about teaching their children to be more polite. Indeed, from very early on, parents teach children to follow normative rituals to say "please", "thank you", "hello" and "good-bye" (Gleason et al., 1984). It can be enlightening to know whether and when children understand positive implications of following these norms.

Third, examining children's *comprehension* of polite speech as compared to their *production* is meaningful, in that children's comprehension can reveal more abstract representations and inferences about language than their productivity (e.g., Fisher, 2002): Children's ability to say "please" early on does not necessarily indicate that they understand saying "please" is more polite, nicer and socially apt, as children may simply obey or imitate what their caregivers tell them to say without understanding its meaning.

Research on children's comprehension of polite speech has received less focus than research on their production of polite speech. Moreover, the few studies that did examine children's understanding of polite speech have been largely inconclusive. Though there was some initial evidence to suggest that producing a request with "please" is judged to be polite by three years of age (E. Bates, 1976; E. Bates & Silvern, 1977), in a later study, the judgment of "please" as being polite was only replicated starting at five years of age, but not younger (Nippold et al., 1982). These initial studies also lacked statistical tests to assess each age group's performance, and did not systematically manipulate cues other than linguistic markers (e.g., prosody or facial expressions).

In addition to children's recognition of politeness markers, there are also many open questions about their abilities to recognize the intentions underlying polite speech. For example, do children know that the word "polite" should be associated with politeness rules people abide by (e.g., saying "please")? Relatedly, do children recognize polite speech as being positively valenced, such that they

think it is better and nicer to say polite things? Do children understand the social implications of speaking politely? For example, polite people may be more likely to get their wishes granted (“I will pour him more water because he was nice”) and may be better social play partners compared to those who are impolite. Finally, what cues to politeness do children recognize? Do they recognize linguistic politeness markers such as “please,” or “can you,” or both? Or do they rely on prosodic cues that make utterances sound more respectful, or on facial expressions that make a person look kind?

In this current work, we sought to answer these questions, and test what 2- to 4-year-old children understand about requests using politeness markers. Across three experiments, we presented stories about speakers who decided to speak politely (e.g., “Please pour me more water”) or impolitely (“Pour me more water”) and asked child participants to make judgments between the two speakers. We examined in each experiment whether: (1) children are able to reason about speakers using polite speech as being relatively more “polite” and “nice” and less “rude” or “mean” than speakers not using polite speech; (2) they can reason about social implications of using polite speech (e.g., politeness as a sign of a nice play partner, or greater likelihood of compliance from the addressee); and (3) they show improvement with age for these lines of reasoning. We also examined whether children need additional cues to politeness such as facial expressions (Expt 1) or prosodic cues (Expt 2), or they can make use of linguistic politeness markers alone (Expt 3) to make appropriate inferences about speakers.

3.2 Experiment 1

In Experiment 1, we tested whether 3- to 4-year-old children were able to understand the implications of using simple politeness markers, based on linguistic cues of interest (whether the speaker says “please,” “can you”) and other cues (facial expressions and prosodic cues) that make polite speech more salient and naturalistic.

3.2.1 Methods

Participants

3-year-old ($n = 20$; 12 F, $M_{age} = 3.61$ years, $SD_{age} = 0.22$) and 4-year-old children ($n = 18$; 6 F, $M_{age} = 4.38$ years, $SD_{age} = 0.25$) were recruited from a local preschool. An additional 3 children were tested but excluded due to failure on the practice questions ($n = 2$) or completion of fewer than half of the test trials ($n = 1$).

Stimuli and design

We designed a picture book with twelve stories in which a protagonist is approached by two speakers, one of whom makes a request by producing an utterance with a politeness marker (e.g., “Please pour me more water”), and the other produces an utterance without (“Pour me more water”). There were three types of politeness marker that could be used: “please” (as in “Please pour me more water”), “can you” (“Can you pour me more water”), and “can you please” (“Can you please pour me more water”).

We designed six question types to ask participants following the presentation of the stories: four *speaker attribute* questions (*polite*: “Which one was more polite?”; *rude*: “Which one was more rude?”; *nice*: “Which one was nicer?”; *mean*: “Which one was meaner?”) and two *social implication* questions (*play partner*: “Which one would you rather play with?”; *compliance*: “Which one will [get what they want]?”). Each participant would be asked one of the four speaker attribute questions, followed by one of the two social implication questions.

In Experiment 1, all utterances were produced live by the experimenter, with appropriate prosodic cues and facial expressions for each request: Utterances with politeness markers were produced by kind voice and facial expression, whereas utterances lacking politeness marker were produced with angry voice and facial cues.

Procedure

The experimenter presented to the child a storybook with a total of thirteen stories about different characters. In the *practice* phase, the child heard a story with one clearly mean character (*Drew kicked Carol*) and one clearly nice character (*Graham gave Carol a gift*). After a reminder of what each character did, the experimenter asked the participant: *Which one was being meaner?* and

Which one was being nicer? If the child answered the question wrong the first time, the experimenter read the story one more time, saying, "Let's think about the story one more time." Only children who correctly answered both questions in the first or second attempt were included in the analyses.

In the *test* phase, the child heard twelve stories, in each of which they saw one speaker who decided to speak politely (*Jean wanted more water in her cup. Jean said to Fred, "Please pour me more water"*) and another speaker who spoke impolitely (*Suzy also wanted more water in her cup. Suzy said to Fred, "Pour me more water."*). After a reminder about what each speaker said, the child was asked a total of two questions. For the first question, the experimenter asked one out of four possible questions for speaker attribute: "Which one was being more polite [more rude/nicer/meaner]?" For the second, social implication question, the experimenter either asked about play partner (*Which one would you rather play with?*) or likelihood of compliance (e.g., *Which one will Fred give water to?*). The order of story types and question types was counterbalanced.

3.2.2 Results and Discussion

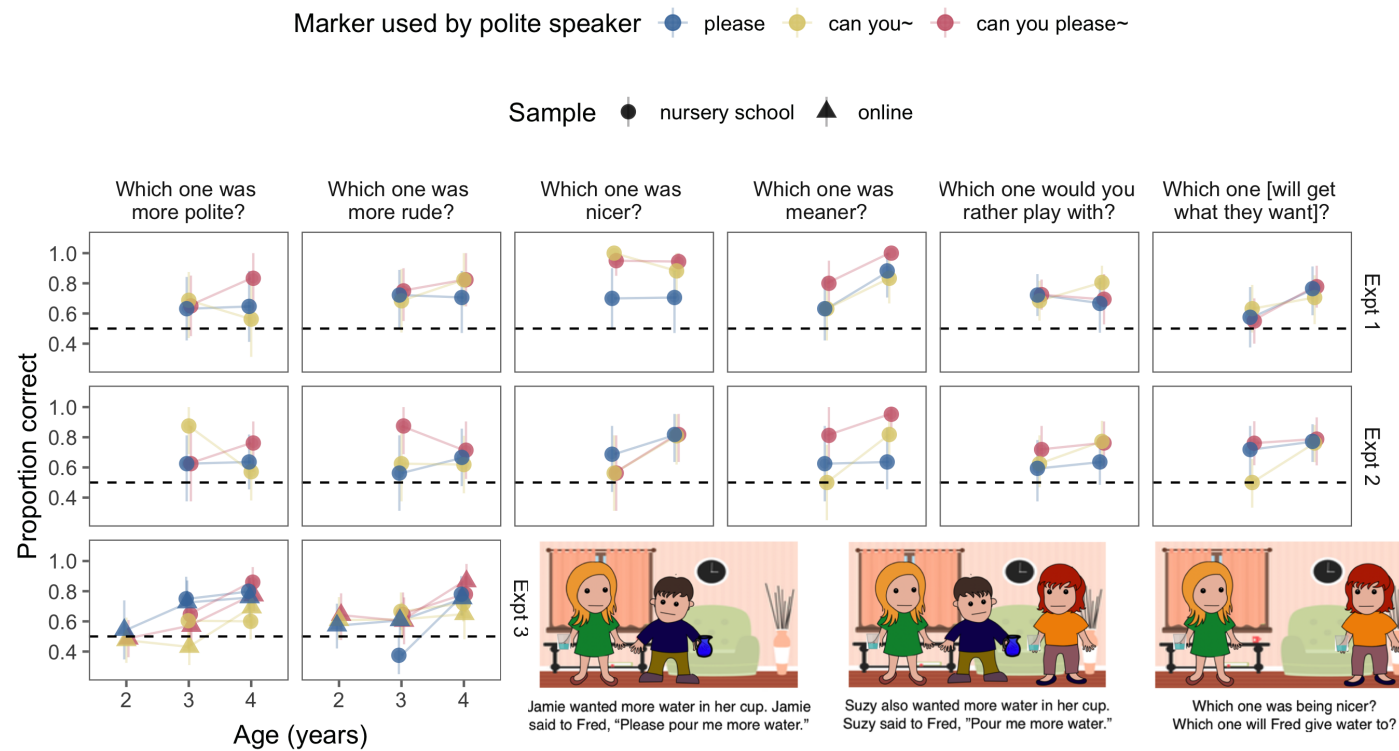


Figure 3.1: Bottom right: Story example. Top, left: Results. Proportion of correct responses to questions comparing between a speaker who used a politeness marker (where blue indicates "please", yellow "can you", and red "can you please") versus a speaker who did not. Data are binned into one-year age groups. Each row represents data from a different Experiment. Columns represent different questions asked. Dashed line represents chance level (i.e., if participant were guessing at random).

We looked at the proportion of correct responses to various questions comparing speakers who used a politeness marker and spoke kindly, and speakers who did not use a politeness marker and spoke meanly (Figure 3.1, first row). A mixed-effects logistic regression predicting accuracy based on age, question type and politeness marker type² showed there was an improvement with age ($\beta = 0.2$, $p = 0.026$). The regression model also revealed that children seemed to find some question types easier than others: Responses to *nice* and *mean* questions were more accurate than to *polite* and *rude* questions ($\beta = 0.8$, $p = 0.002$), whereas social implication questions (*play partner* and *compliance*) were overall more difficult compared to speaker attribute questions (*polite*, *rude*, *nice*, and *mean*; $\beta = -0.33$, $p = 0.006$).

Looking more closely at responses for each of the question types, children from both age groups tended to accurately answer the *polite*, *nice*, *mean*, *rude*, and *play partner* questions overall (3-year-olds' mean accuracy range: 0.58 - 0.88; 4-year-olds' mean accuracy range: 0.68 - 0.9), indicating correctly that the speaker who used a politeness marker was more polite and nicer, and less mean and rude, and was likely a better play partner. For the *compliance* question, 4-year-olds overall answered correctly that the speaker who used politeness marker will likely get what they want from the listener ($M_{4y} = 0.75$, $p < .01$), but 3-year-olds did not perform above chance ($M_{3y} = 0.58$). As for the different politeness marker types, both age groups overall tended to give correct answers based on all three markers, but especially "can you please" (3-year-olds: $M_{please} = 0.66$, $M_{canyou} = 0.72$, $M_{canyouplease} = 0.74$; 4-year-olds: $M_{please} = 0.73$, $M_{canyou} = 0.77$, $M_{canyouplease} = 0.84$).

In sum, in this first experiment, we saw preliminary evidence that children pay attention to some cues to politeness and are able to use these cues to infer whether speakers are relatively polite, rude, nice or mean, and whether speakers are good play partners and are likely to get what they wanted from their addressees. 4-year-olds answered questions accurately more often compared to 3-year-olds, especially for the question about addressee's compliance with the speaker's request. In general, however, both age groups tended to be accurate when all the possible cues were used to signal that one speaker was polite (used "can you please", spoke with a kind tone and face) and the other speaker wasn't (did not use a politeness marker, spoke with an angry tone and face).

There were a number of remaining issues from Experiment 1. Children may not have used

²for Experiments 1 and 2, we use this model structure with a maximal random effect structure that converges: `accuracy ~ age x question type x politeness marker type + (1 | item)`, where age is continuous, centered and scaled. All categorical variables were deviation coded, with specified contrasts of interest for the question type. Significance was calculated using the standard normal approximation to the t distribution (Barr, Levy, Scheepers, & Tily, 2013a).

the linguistic politeness markers (e.g., “please”) per se, and rather prosodic and facial cues that accompany these markers. That is, children may have relied on the speaker’s kind voice and face rather than their use of “please” to evaluate their niceness or likeability as a play partner. Similarly, greater accuracy for some questions over others (e.g., *nice* > *polite*) may have been due to greater association between some of the words and prosodic and facial cues (e.g., a kind face may be seen to signal niceness more than politeness), not due to greater understanding for those words or concepts. Another concern is that the experimenter was aware of the manipulations (i.e., they knew which speaker was supposed to be “polite”) and thus could have affected the presentation of these speakers in ways that are not consistent across all participants. In our next two experiments, we sought to remove these potential confounds.

3.3 Experiment 2

In Experiment 1, we saw initial evidence that children can use some combinations of linguistic, prosodic, and facial cues to politeness. In Experiment 2, we examined whether children can make similar judgments using linguistic and prosodic cues only, without facial expressions. For this, we conducted a preregistered experiment where we used pre-recorded voiceovers to present speaker utterances, so that (1) we could look at children’s judgments based on linguistic markers and prosodic cues only, and (2) we could remove the role of the experimenter in presentation of these utterances.

3.3.1 Methods

Participants

3-year-old ($n = 16$; 8 F, $M_{age} = 3.56$ years, $SD_{age} = 0.29$) and 4-year-old children ($n = 22$; 13 F, $M_{age} = 4.5$ years, $SD_{age} = 0.32$) were recruited from a local preschool. An additional 5 children were tested but excluded due to failure on the practice questions.

Stimuli and design

The design was identical to Experiment 1. Stimuli were the same as Experiment 1 except two changes: (1) Instead of a picture book, we presented the stories on a tablet; (2) the speakers’ utterances were now presented as recorded voiceovers. The voiceovers were recorded by native

English speakers, and contained prosodic cues that matched the presence/absence of a politeness marker (e.g., “Please pour me more water” was recorded with a kind voice and “pour me more water” with an angry voice).

Procedure

The procedure was identical to Experiment 1, except for the following change: The participants now had to tap on a speaker on tablet in order either to hear them speak, or to choose an answer to the questions asked.

3.3.2 Results and Discussion

Overall we saw similar patterns of results in Experiment 2 (Figure 3.1, second row) compared to Exp. 1. A mixed-effects logistic regression predicting accuracy based on age, question type and politeness marker type showed that accuracy improved with age ($\beta = 0.25$, $p = 0.002$), and children made accurate judgments more often when the politeness marker was “can you please” than when the marker was “please” or “can you” ($\beta = 0.33$, $p = 0.019$). There was no main effect of question type, but there was an interaction between age and question type such that performance for *nice* and *mean* questions saw greater improvement with age than for *polite* and *rude* questions ($\beta = 0.57$, $p = 0.011$).

For children’s responses to different question types, 3-year-olds’ accuracy did not differ from chance level for *nice*, *mean*, and *play partner* questions, but their means numerically exceeded 50% for all question types, and 4-year-olds accurately answered questions of all types (3-year-olds’ mean accuracy range: 0.6 - 0.88; 4-year-olds’ mean accuracy range: 0.66 - 0.9). For politeness marker types, 3-year-olds’ performance did not differ from chance for “please” and “can you”, but both age groups tended to answer questions about different politeness markers accurately overall (3-year-olds: $M_{please} = 0.63$, $M_{canyou} = 0.61$, $M_{canyouplease} = 0.72$; 4-year-olds: $M_{please} = 0.7$, $M_{canyou} = 0.72$, $M_{canyouplease} = 0.8$).

In sum, across Experiments 1 and 2, we saw that children tend to make accurate judgments about speakers given their use of politeness markers, especially “can you please,” together with prosodic cues, and children get better with age in their use of politeness cues to respond to questions about speaker attributes and social implications.

3.4 Experiment 3

We conducted a third, pre-registered experiment to see whether children are able to evaluate speakers based on linguistic markers only, without any other supporting cues such as prosodic cues or facial expressions.

3.4.1 Methods

Participants

We recruited two samples of participants: one from the same local nursery school as Experiments 1 and 2, and the other from Lookit (<https://lookit.mit.edu/>), an online platform for child research participation, in which parents and their children can participate together. The nursery school sample consisted of 3-year-old ($n = 24$; 11 F, $M_{age} = 3.65$ years, $SD_{age} = 0.26$) and 4-year-old children ($n = 25$; 13 F, $M_{age} = 4.48$ years, $SD_{age} = 0.28$). An additional 3 children were tested but excluded due to failure on the practice questions. The online sample consisted of 2-year-old ($n = 23$; 12 F, $M_{age} = 2.48$ years, $SD_{age} = 0.29$), 3-year-old ($n = 31$; 15 F, $M_{age} = 3.59$ years, $SD_{age} = 0.27$) and 4-year-old children ($n = 27$; 12 F, $M_{age} = 4.46$ years, $SD_{age} = 0.29$). An additional 28 children were tested but excluded due to failure on the practice questions ($n = 19$) or completion of fewer than half of the test trials ($n = 9$).

Stimuli

For the nursery school sample, stimuli were identical to Experiment 2 except that the voiceovers for all utterances had the same prosody: All utterances ended with a rising intonation. For the online sample, stimuli were identical to what the nursery school participants saw except that the story narrations (other than speaker utterances) were also pre-recorded such that parents did not need to read the stories aloud to their children.

Procedure

For the nursery school sample, the procedure was identical to Experiment 2. For the online sample, the procedure was similar except that parents and children participated together at home and there was no experimenter present. Parents accessed the webpage for the study and gave their consent for

participation, and then read instructions to proceed through the different stories, which specifically asked the parents to not tell their children correct answers for the questions.

3.4.2 Results and Discussion

Experiment 3

For Experiment 3, we were able to look at how children answered the *polite* and *rude* questions given the same three politeness marker types as in Experiments 1 and 2, with three age groups including 2-year-olds. (Fig. 3.1, third row).

A mixed-effects logistic regression controlling for the effect of sample³ showed improvement with age ($\beta = 0.19$, $p = 0.033$) as well as better performance for “can you please” than “please” and “can you” together ($\beta = 0.42$, $p = 0.002$), consistent with Experiment 2 results. Performance for “please” was also better than for “can you please” and “please” together ($\beta = 0.3$, $p = 0.027$), which may be surprising given that we previously did not see the same effect in Experiments 1 and 2. One possible explanation is that controlling for prosodic cues in Experiment 3 actually made it *easier* to use “please” as a politeness cue. Because we had stripped all the other variations, it may have made the contrast between the presence and absence of the marker “please” *more* salient.

Additionally, children were better with the *polite* questions than *rude* overall ($\beta = -0.19$, $p = 0.04$), but especially given “please” ($\beta = 0.42$, $p = 0.002$). Finally, children showed a greater improvement with age for “can you please” compared to “please” and “can you” together ($\beta = 0.38$, $p = 0.004$).

All experiments

Did children perform better given facial and/or prosodic cues, or were linguistic politeness markers sufficient? To see any potential effect of experiment on children's performance, we conducted an exploratory mixed-effects logistic regression on all three experiments together⁴. The regression model showed no significant main effect of experiment, suggesting that children did not perform more poorly when facial and prosodic cues were removed, and they were able to make accurate judgments based on linguistic cues alone. The model also showed that children improved with increasing age ($\beta =$

³Model structure: `accuracy ~ sample + age x question type x politeness marker type + (1 | item)`

⁴Model structure: `accuracy ~ sample + experiment + age x question type x politeness marker type + (1 | item)`

0.33, $p < .001$) and that children were more accurate with “can you please” than “please” and “can you” ($\beta = 0.25$, $p = 0.011$), confirming results from each individual experiment. Additionally, the model showed that children became better at judging the politeness marker “can you please” with age ($\beta = 0.73$, $p = 0.005$), and that children answered *polite* questions better than *rude* questions about the marker “please” ($\beta = 0.26$, $p = 0.006$)

3.5 General Discussion

What do young children understand about polite speech? In three experiments, we looked at how 2- to 4-year-old children reason about making requests with or without simple politeness markers such as “please”, “can you” and “can you please.” By 3 years, children pay attention to the use of politeness markers to accurately judge whether that speaker is relatively more polite, rude, nicer or meaner compared to another speaker. By 4 years, children reliably infer that a speaker who uses a politeness marker is a better play partner and more likely to get what they want. Across all three experiments, we saw a clear developmental trend such that children improved in their reasoning about polite speech with increasing age. We observed no large experiment effects as we eliminated facial and prosodic cues; instead, all these inferences appeared to be supported by linguistic markers alone.

Even though children have been shown to produce polite speech such as “please,” evidence has been sparse and inconclusive for whether young children below 5 years comprehend speaker attributes and intentions based on polite speech. Here, we found that children are sensitive to the use of politeness markers in speech, and are able to use these markers to infer the speaker’s attributes (e.g., niceness) by 3 years, and consequent social implications by 4 years. These ages are closer to the age of first reliable production of polite speech than have been suggested by earlier work.

Children in the US are often explicitly taught and prompted to use politeness markers such as “please” in their requests from early on (e.g., “What’s the magic word?”; Gleason et al., 1984), thus they may quickly learn to use these markers as a rule in order to get what they want. They also might hear other remarks that pair politeness markers with positive words (e.g., “You should be *nice* and say *please*”), which may help them learn the association between polite speech and positive attributes. Gradually, children may recognize more subtle social processes that are related to polite speech production: Adults may praise and reward children who spoke politely, and children

themselves may like peers who ask for permission to play with their toys rather than take the toys away without asking. Future work with corpus data analysis looking at these interactions between children and others may reveal important conversational patterns that help children acquire social meanings of polite speech.

There are limitations to the current work that present other opportunities for future research. Because this work looked only at the behaviors of English-speaking children with a relatively high socioeconomic status in the US, it is an open question how children with different language and cultural background may develop understanding of polite speech. Cross-cultural investigation of what markers are present in other languages, cultures and backgrounds, as well as how those markers are acquired, will be informative.

Also, we did not manipulate the social status of speakers or addressees. Though not explicitly stated, the visual depiction and narration used for the current work suggested that speakers were communicating with their peers only. However, one key prediction from politeness theory is that speakers will adjust their utterances based on the status of the addressees (P. Brown & Levinson, 1987). Indeed children adjust own their speech based on the listener status and age: Even at two years, children use a polite form of request (“Can I have...”) to an adult but an imperative form (“Give me...”) to a peer (Shatz & Gelman, 1973). Thus, future work should examine how children use cues to politeness to judge speaker intentions in different contexts, including varied status differences between speakers and listeners.

In sum, the current work showed that young children understand implications of using simple politeness markers in requests. A broader understanding of the emergence of politeness may offer insights into how children become proficient users of language across the wide range of social situations that they encounter.

Chapter 4

Children understand tradeoffs between informational and prosocial goals

FIXME.

FIXME.

Conclusion

In this dissertation, I proposed a framework for understanding children’s information-seeking decisions within social contexts. The core of the argument is that the presence of other people can change the *availability* and *usefulness* of information-seeking behaviors by shaping learners’ goals, hypotheses, actions, answers, and thresholds for stopping information gathering. Following the theoretical framework, I presented a set of empirical studies that explored whether the dynamics of children’s real-time information selection via their eye movements flexibly adapts to gather social information that supports language processing.

Chapter 2 investigated how children learning American Sign Language (ASL) distributed visual attention between the linguistic signal and referents, which both compete for visual attention. Similar to children learning spoken language, ASL learners shifted gaze away from a social partner to seek objects before sign offset, providing evidence that, despite channel competition, language drove rapid shifts in visual attention to named referents. Chapter 3 extended the sign language research by directly comparing ASL learners’ gaze dynamics to those of children learning spoken English using parallel real-time language comprehension tasks. Chapter 3 also presented a comparison of English-learning children and adults’ eye movements in noisy vs. clear auditory contexts. In both the ASL and noisy speech cases, listeners adapted their gaze to seek additional language-relevant information from social partners before shifting to seek a named referent. Chapters 4 and 5 explored how eye movements change when children and adults processed familiar and novel words accompanied by social cues to reference. Taken together, the social gaze work suggests that children integrate their uncertainty over word-object mappings to decide when to seek social information, which in turn, modulates the input to statistical word learning.

The integrative framework and empirical work described here are limited in several important

ways. First, the majority of this research tested binary hypotheses of behavior change – i.e., sign vs. spoken language; noisy vs. clear speech; word learning with vs. without social gaze – to answer the question of whether children would flexibly adapt their information seeking in response to changes in their processing environments. Chapters 2-5 present evidence across a diverse set of case studies that children's real-time information seeking is quite flexible. However, to move the integrative framework forward, we would want to develop a fully-specified model that could make quantitative predictions about how social contexts will change the utility of information-seeking behaviors. This step will require formalizing the notions of value and cost of information-seeking actions in a modeling framework that can incorporate the effects of reasoning about other people's mental states.

We have taken some initial steps towards this goal by developing a model of active-social learning that integrates ideas from Optimal Experiment Design (OED) with formalizations of recursive social reasoning from Bayesian models of pragmatic language interpretation (N. D. Goodman & Frank, 2016b). We found that this integrated model was able to reproduce the qualitative patterns in adults' decisions of whether to forego information seeking in favor of more immediately rewarding actions when their social partner highlighted performance and presentational goals (Erica J Yoon, MacDonald, Asaba, Gweon, & Frank, 2018). The integrative framework described here directly inspired this line of research, and I hope that future versions of the model will be able to generate graded, testable predictions for behavior across a variety of domains – e.g., eye movements, early vocalizations, and verbal question asking.

Second, we used one particular formalization of active inquiry. The OED model focuses on learners' information-seeking decisions given a specific goal to learn and a set of candidate hypotheses. Other computational frameworks have formalized active learning in different ways. For example, foraging models pursue the analogy that human information seeking is similar to animals' decisions about where and how long to look for food if they were trying to maximize caloric intake while minimizing their effort and time (see Pirolli & Card (1999) for a review). Cognitive scientists have successfully modeled a range of behaviors as a form of spatial foraging, such as searching for semantic concepts in memory (Hills, Jones, & Todd, 2012) and decisions about where to direct visual attention in real-time (Manohar & Husain, 2013). In addition to these search models, recent work on curiosity-based learning in developmental robotics has created algorithms that optimize intrinsic estimates of learning progress. This formalization creates systems that focus on seeking

activities and stimuli of intermediate complexity where learners' predictions are steadily improving, and uncertainty is steadily decreasing (Oudeyer & Smith, 2016). One of the challenges for researchers trying to integrate active and social learning is that the space of possible connections is quite large. By constraining our framework to active decision making, we were able to make some progress on an important sub-component of a larger set of children's information-seeking behaviors. Future theoretical work, however, should consider possible connections between social learning phenomena and the foraging/curiosity-based learning frameworks.

Third, our ultimate goal for the active-social framework is to incorporate effects at a developmental timescale. The experiments in this thesis, however, often treated children and adults as two endpoints on a continuum, exploring parallels and differences between children's information seeking and our best estimate of the mature state of the language processing system. We did find some clear patterns of developmental change. In Chapter 3, adults were faster to respond to familiar words, generated more language-consistent shifts, and produced fewer early shifts before accumulating enough information. Older children were also faster to respond than younger children but did not generate more language-consistent shifts overall. Older children did, however, produce fewer early, non-language-driven gaze shifts. This pattern of results suggests that what might be developing is an ability to inhibit a behavior – shifting gaze away from the language source – that reduces access to information that is useful for figuring out the identity of a named referent. Prior work also shows that children develop greater flexibility in ignoring irrelevant information to focus on parts of the meaningful parts of a sentence (Zangl & Fernald, 2007). But it is still an open question as to how children's visual information seeking might change as they become more efficient in processing words and develop their skill in focusing on relevant information in their environment.

In addition to change at the developmental timescale, the final experiment in Chapter 5 represents an exception where we measured adaptation of information seeking over the slightly longer timescale of multiple exposures to novel word-object links and in the context of highly familiar words, which children learned through exposure to many prior labeling events in their day-to-day lives. While the study in Chapter 5 is a useful first step, future work should measure change over a longer timescale by densely sampling children at different ages and points of development. For example, it would be useful to know the effect of children's rapidly improving productive language skills, which increases the set of information-seeking actions available by allowing children to ask verbal questions. One

prediction based on our framework is that seeking social information via eye movements should become less useful when children can produce the verbal question “What is this thing called?” since it has a higher probability of returning useful information. Another example is children’s rapid theory of mind development. Our framework predicts that young children should focus more on learning goals if they are less skilled at reasoning about others’ beliefs. But, as their social reasoning abilities mature and their social environments become more complex, children may forego information seeking actions that make them appear incompetent to their social partners.

Fourth, our account is currently underdeveloped concerning individual differences. That is, the model was designed to explore general principles about how qualitative changes to the social environment might shape children’s information seeking actions. However, it is possible to use the active-social framework to understand how individual differences in children’s input and cognitive abilities might interact to shape how they decide to seek information from social partners. For example, prior research has found that adults vary in the proportion of unambiguous naming episodes they provide, with some parents rarely providing highly informative contexts and others’ doing so relatively more often (Cartmill et al., 2013). Within our active-social framework, this differential experience could be instantiated as children learning a model of the probability of getting a high-quality answer when they ask a question. If children do not expect an answer is likely, then this should reduce information seeking even if there is a social partner available. We did find some evidence of this effect in Chapters 4 and 5 where adults were less likely to use an unreliable social cue to reference and where children looked more to a social partner who provided gaze cue than to one who did not.

Individual differences in cognitive abilities could also be included in our model. Prior research shows variability in children’s theory of mind and inhibitory control abilities (Carlson & Moses, 2001), in addition to the considerable variability in language processing skill (Marchman & Fernald, 2008). Within our active-social framework, children with a more-developed theory of mind skill might place a higher weight on pursuing social goals over and above informational goals, taking actions that maintain others’ beliefs about their abilities. It could also be that stronger perspective-taking skills help children reason about the probability of getting a quality answer from another person, thus modulating whom they choose to ask questions (e.g., seeking information from an

expert vs. a novice). Another compelling possibility is that children who have stronger domain-general processing abilities are better able to update their beliefs based on the information they receive, and thus reducing the amount of time they spend seeking information from social partners. These are all interesting, open questions for future research that fall out of the integrative active-social framework proposed in this thesis.

Finally, the empirical research described here aimed to understand how children's information-seeking behaviors adapt to support their language processing. To accomplish this, we measured changes in children's gaze dynamics during language comprehension and word learning in simplified environments. This approach has the benefit of providing a high degree of experimenter control and a relatively well-understood hypothesis linking observable behavior (eye movements) to underlying psychological constructs (e.g., lexical access) (Tanenhaus, Magnuson, Dahan, & Chambers, 2000). The risk, however, is that the responses that we can measure in the lab do not reflect behaviors that support children's learning in their natural environments. That is, children acquire their first language from conversations where social partners produce contingent responses and take actions that control the flow of children's learning experience. This gap suggests two critical next steps for the research described here: (1) measure changes in children's information seeking within free-flowing social interactions with their caregivers (see Franchak, Kretch, & Adolph (2018) for a recent example of this approach using head-mounted eye trackers), and (2) develop more realistic lab-based experiments that incorporate behaviorally-relevant features of children's learning environments such as contingent responding to children's actions (see Benitez & Saffran (2018) for an example of studying word learning using a gaze-contingent eye-tracking paradigm).

In sum, we set out to explore how children's eye movements adapt to a wide range of social contexts during two ecologically-relevant tasks: familiar language comprehension and novel word learning. We found that children could adapt their gaze to seek relevant social information when it was useful for language processing. Moreover, children and adults showed evidence of differential learning of new words when social gaze guided their visual attention. This work highlighted two critical, open challenges for a framework of information-seeking within social contexts: (1) develop a precise quantitative model of how social learning can change the utility of information-seeking behaviors, and (2) move beyond highly-constrained lab experiments to document information seeking behaviors in children's natural learning environments. Despite these challenges, the integrative

framework presented in this thesis represents a way forward for understanding how children's information seeking adapts to the wide variety of social environments in which children acquire their first language.

Appendix A

Supplementary materials for Chapter 2

In this appendix, we present FIXME.

A.1 Model details

The *literal listener* L_0 is a simple Bayesian agent that takes the utterance to be true:

$$P_{L_0}(s|w) \propto \llbracket w \rrbracket(s) * P(s).$$

where $\llbracket w \rrbracket(s)$ is the truth-functional denotation of the utterance w (i.e. the utterance’s literal meaning): It is a function that maps world-states s to Boolean truth values. The literal meaning is used to update the literal listener’s prior beliefs over world states $P(s)$.

The *speaker* S_1 chooses utterances approximately optimally given a utility function, which can be decomposed into two components. First, informational utility (U_{inf}) is the amount of information a literal listener L_0 would still not know about world state s after hearing a speaker’s utterance w . Second, social utility (U_{soc}) is the expected subjective utility of the state inferred given the utterance w . The utility of an utterance subtracts the cost $c(w)$ from the weighted combination of the social and epistemic utilities.

$$U(w; s; \phi_{S_1}) = \phi_{S_1} \cdot \ln(P_{L_0}(s \mid w)) + (1 - \phi_{S_1}) \cdot \mathbb{E}_{P_{L_0}(s|w)}[V(s)] - C(w).$$

The speaker then chooses utterances w softmax-optimally given the state s and his goal weight mixture ϕ_{S_1} :

$$P_{S_1}(w \mid s, \phi_{S_1}) \propto \exp(\lambda_1 \cdot \mathbb{E}[U(w; s; \phi_{S_1})]).$$

A.2 Literal semantic task

We probed judgments of literal meanings of the target words assumed by our model and used in our main experiment.

A.2.1 Participants

51 participants with IP addresses in the United States were recruited on Amazon’s Mechanical Turk.

A.2.2 Design and Methods

We used thirteen different context items in which a speaker evaluated a performance of some kind. For example, in one of the contexts, Ann saw a presentation, and Ann’s feelings toward the presentation (true state) were shown on a scale from zero to three hearts (e.g., two out of three hearts filled in red color; see Figure 2.3 for an example of the heart scale). The question of interest was “Do you think Ann thought the presentation was / wasn’t X?” and participants responded by choosing either “no” or “yes.” The target could be one of four possible words: *terrible*, *bad*, *good*, and *amazing*, giving rise to eight different possible utterances (with negation or no negation). Each participant read 32 scenarios, depicting every possible combination of states and utterances. The order of context items was randomized, and there were a maximum of four repeats of each context item per participant.

A.2.3 Behavioral results

We analyzed the data by collapsing across context items. For each utterance-state pair, we computed the posterior distribution over the semantic weight (i.e., how consistent X utterance is with Y state)

assuming a uniform prior over the weight (i.e., a standard Beta-Binomial model). Meanings of the words as judged by participants were as one would expect (Figure A.1).

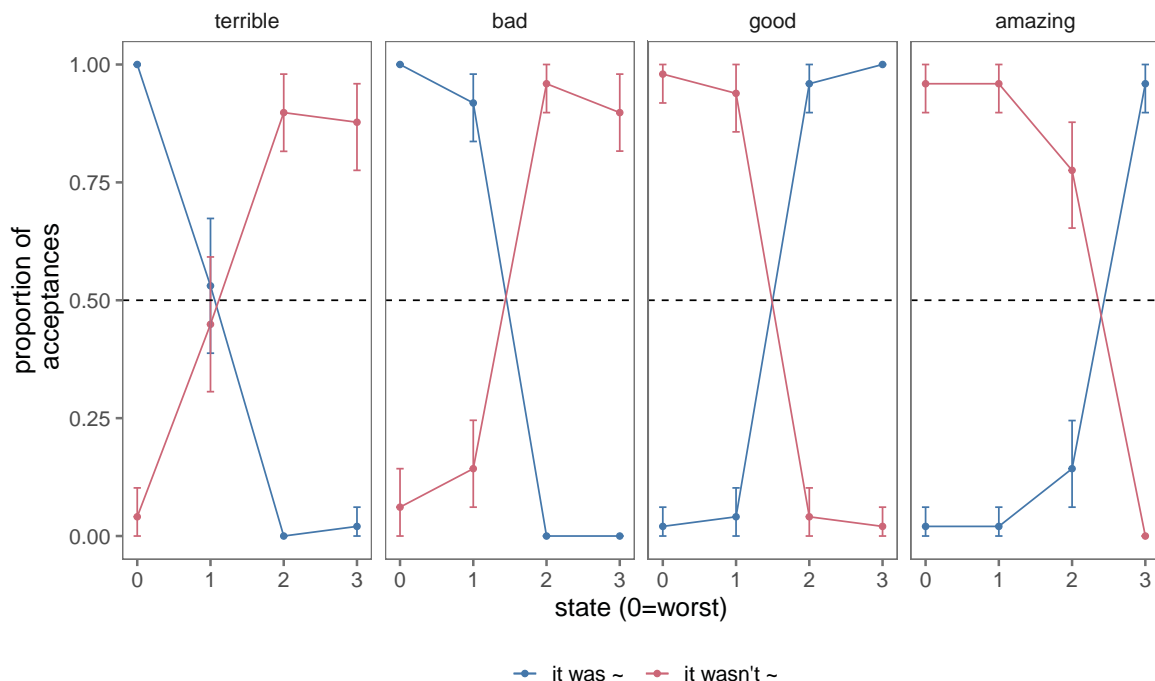


Figure A.1: Semantic measurement results. Proportion of acceptances of utterance types (shown in different colors) combined with target words (shown in different facets) given the true state represented on a scale of hearts. Error bars represent 95% confidence intervals.

A.3 Data analysis

We used R (Version 3.4.3; R Core Team, 2017) and the R-packages *BayesFactor* (Version 0.9.12.2; Morey & Rouder, 2015), *bindrcpp* (Version 0.2.2; MÅijller, 2017a), *binom* (Version 1.1.1; Dorai-Raj, 2014), *brms* (Version 2.0.1; BÅijrkner, 2017), *coda* (Version 0.19.1; Plummer, Best, Cowles, & Vines, 2006), *directlabels* (Version 2017.3.31; Hocking, 2017), *dplyr* (Version 0.8.0.1; Wickham, Francois, Henry, & MÅijller, 2017), *forcats* (Version 0.2.0; Wickham, 2017a), *ggplot2* (Version 3.0.0; Wickham, 2009), *ggthemes* (Version 3.4.0; Arnold, 2017), *gridExtra* (Version 2.3; Auguie, 2017), *here* (Version 0.1; MÅijller, 2017b), *jsonlite* (Version 1.6; Ooms, 2014), *langcog* (Version 0.1.9001; Braginsky, Yurovsky, & Frank, n.d.), *lme4* (Version 1.1.15; D. Bates, MÅdchler, Bolker, & Walker, 2015),

Table A.1: Predictor mean estimates with standard deviation and 95% credible interval information for a Bayesian linear mixed-effects model predicting negation production based on true state and speaker goal (with both-goal as the reference level).

Predictor	Mean	SD	95% CI-Lower	95% CI-Upper
Intercept	0.88	0.13	0.63	1.12
True state	2.18	0.17	1.86	2.53
Goal: Informative	0.47	0.17	0.14	0.80
Goal: Kind	0.97	0.25	0.51	1.49
True state * Informative	-1.33	0.18	-1.69	-0.98
True state * Kind	-0.50	0.22	-0.92	-0.07

magrittr (Version 1.5; Bache & Wickham, 2014), *Matrix* (Version 1.2.12; D. Bates & Maechler, 2017), *papaja* (Version 0.1.0.9655; Aust & Barth, 2017), *purrr* (Version 0.2.5; Henry & Wickham, 2017), *RColorBrewer* (Version 1.1.2; Neuwirth, 2014), *Rcpp* (Eddelbuettel & Balamuta, 2017; Version 1.0.1; Eddelbuettel & François, 2011), *readr* (Version 1.1.1; Wickham, Hester, & François, 2017), *rwebpl* (Version 0.1.97; Braginsky, Tessler, & Hawkins, n.d.), *stringr* (Version 1.3.1; Wickham, 2017b), *tibble* (Version 2.1.1; Mäijler & Wickham, 2017), *tidyr* (Version 0.7.2; Wickham & Henry, 2017), and *tidyverse* (Version 1.2.1; Wickham, 2017c) for all our analyses.

A.4 Full statistics on human data

We used Bayesian linear mixed-effects models (*brms* package in R; Bürkner, 2017) using crossed random effects of true state and goal with maximal random effects structure (Barr, Levy, Scheepers, & Tily, 2013b; A. Gelman & Hill, 2006). The full statistics are shown in Table A.1.

A.5 Model fitting and inferred parameters

Other than speaker goal mixture weights explained in the main text (shown in Table 2.2), the full model has two global parameters: the speaker’s soft-max parameter λ_{S_2} and soft-max parameter of the hypothetical speaker that the pragmatic listener reasons about λ_{S_1} . λ_{S_1} was 1, and λ_{S_2} was inferred from the data: We put a prior that was consistent with those used for similar models in this model class: $\lambda_{S_2} \sim \text{Uniform}(0, 20)$. Finally, we incorporate the literal semantics data into the RSA

Table A.2: Inferred negation cost and speaker optimality parameters for all model variants.

Model	Cost of negation	Speaker optimality
ninformational only	1.58	8.58
ninformational, presentational	1.89	2.93
ninformational, social	1.11	3.07
ninformational, social, presentational	2.64	4.47
presentational only	2.58	9.58
social only	1.73	7.23
social, presentational	2.49	5.29

model by maintaining uncertainty about the semantic weight of utterance w for state s , for each of the states and utterances, and assuming a Beta-Binomial linking function between these weights and the literal semantics data (see *Literal semantics task* above). We infer the posterior distribution over all of the model parameters and generate model predictions based on this posterior distribution using Bayesian data analysis (M. D. Lee & Wagenmakers, 2014). We ran 4 MCMC chains for 80,000 iterations, discarding the first 40,000 for burnin. The inferred values of parameters are shown in Table A.2.

A.6 Data Availability

Our model, preregistration of hypotheses, procedure, data, and analyses are available at https://github.com/ejyoon/polite_speaker.

A.7 Supplemental Figures

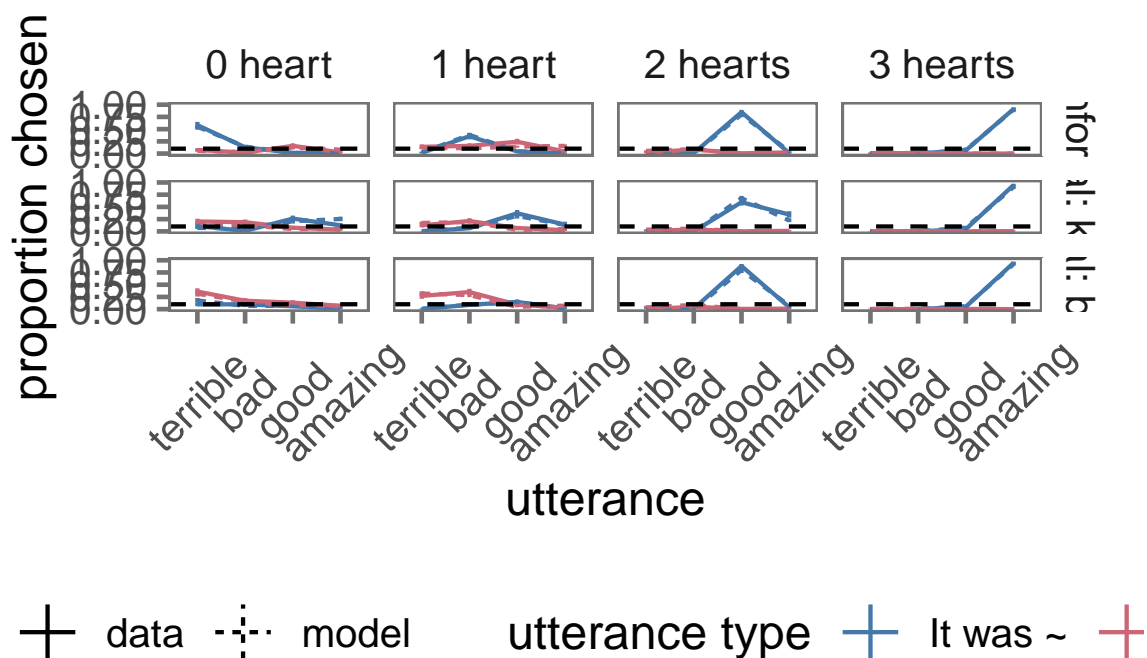


Figure A.2: Experimental results (solid lines) and fitted predictions from the full model (dashed lines) for speaker production. Proportion of utterances chosen (utterance type "direct" vs. "indirect" in different colors and words shown on x-axis) given the true states (columns) and speaker goals (rows). Error bars represent 95% confidence intervals for the data and 95% highest density intervals for the model. Black dotted line represents the chance level.




Figure A.3: Comparison of predictions for proportion of utterances chosen by pragmatic speaker from possible model variants (left) and human data (rightmost) for average proportion of negation produced among all utterances, given true state of 0 heart and speaker with a goal to be informative (top), kind (middle), or both (bottom). Gray dotted line indicates chance level at 12.5%.




Figure A.4: Experimental results (left) and fitted model predictions (right) for average proportion of negation produced among all utterances, given true states (x-axis) and goals (colors).

Appendix B

Supplementary materials for Chapter 3

B.1 Model output for Experiment 3.1

B.2 Model output for Experiment 3.2

References

- Airenti, G., & Angeleri, R. (2011). Situation-sensitive use of insincerity: Pathways to communication in young children. *British Journal of Developmental Psychology*, 29(4), 765–782.
- Angel, E. (2000). *Interactive computer graphics : A top-down approach with opengl*. Boston, MA: Addison Wesley Longman.
- Angel, E. (2001a). *Batch-file computer graphics : A bottom-up approach with quicktime*. Boston, MA: Wesley Addison Longman.
- Angel, E. (2001b). *Test second book by angel*. Boston, MA: Wesley Addison Longman.
- Arnold, J. B. (2017). *Ggthemes: Extra themes, scales and geoms for 'ggplot2'*. Retrieved from <https://CRAN.R-project.org/package=ggthemes>
- Attardo, S. (1997). Locutionary and perlocutionary cooperation: The perlocutionary cooperative principle. *Journal of Pragmatics*, 27(6), 753–779.
- Auguie, B. (2017). *GridExtra: Miscellaneous functions for “grid” graphics*. Retrieved from <https://CRAN.R-project.org/package=gridExtra>
- Augustine, S. (1952). *Treaties on various issues*. Washington, DC: Catholic University of America Press.
- Aust, F., & Barth, M. (2017). *papaja: Create APA manuscripts with R Markdown*. Retrieved from <https://github.com/crsh/papaja>
- Axia, G., & Baroni, M. R. (1985). Linguistic politeness at different age levels. *Child Development*,

918–927.

- Bache, S. M., & Wickham, H. (2014). *Magrittr: A forward-pipe operator for r*. Retrieved from <https://CRAN.R-project.org/package=magrittr>
- Baker, C. L., Jara-Ettinger, J., Saxe, R., & Tenenbaum, J. B. (2017). Rational quantitative attribution of beliefs, desires and percepts in human mentalizing. *Nature Human Behaviour*, 1(4), 0064.
- Baker, C. L., Saxe, R., & Tenenbaum, J. B. (2009). Action understanding as inverse planning. *Cognition*, 113(3), 329–349.
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013a). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278.
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013b). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278.
- Bates, D., & Maechler, M. (2017). *Matrix: Sparse and dense matrix classes and methods*. Retrieved from <https://CRAN.R-project.org/package=Matrix>
- Bates, D., MÄdchler, M., Bolker, B., & Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, 67(1), 1–48. <http://doi.org/10.18637/jss.v067.i01>
- Bates, E. (1976). Acquisition of polite forms: Experimental evidence. *Language and Context: The Acquisition of Pragmatics*, 295–326.
- Bates, E., & Silvern, L. (1977). Social adjustment and politeness in preschoolers. *Journal of Communication*, 27(2), 104–111.
- Baxter, L. A. (1984). An investigation of compliance-gaining as politeness. *Human Communication Research*, 10(3), 427–456.
- Becker, J. A. (1986). Bossy and nice requests: Children’s production and interpretation. *Merrill-Palmer Quarterly (1982-)*, 393–413.
- Benitez, V. L., & Saffran, J. R. (2018). Predictable events enhance word learning in toddlers.

- Current Biology*, 28(17), 2787–2793.
- Bernicot, J. (1991). French children's conception of requesting: The development of metapragmatic knowledge. *International Journal of Behavioral Development*, 14(3), 285–304.
- Bernicot, J., & Legros, S. (1987). Direct and indirect directives: What do young children understand? *Journal of Experimental Child Psychology*, 43(3), 346–358.
- Bernicot, J., Laval, V., & Chaminaud, S. (2007). Nonliteral language forms in children: In what order are they acquired in pragmatics and metapragmatics? *Journal of Pragmatics*, 39(12), 2115–2132.
- Blum-Kulka, S. (1987). Indirectness and politeness in requests: Same or different? *Journal of Pragmatics*, 11(2), 131–146.
- Blum-Kulka, S., Danet, B., & Gherson, R. (1985). The language of requesting in israeli society. In *Language and social situations* (pp. 113–139). Springer.
- Bock, J. K., & Hornsby, M. E. (1981). The development of directives: How children ask and tell. *Journal of Child Language*, 8(01), 151–163.
- Bonnefon, J.-F., & Villejoubert, G. (2006). Tactful or doubtful? Expectations of politeness explain the severity bias in the interpretation of probability phrases. *Psychological Science*, 17(9), 747–751.
- Bonnefon, J.-F., Dahl, E., & Holtgraves, T. M. (2015). Some but not all dispreferred turn markers help to interpret scalar terms in polite contexts. *Thinking & Reasoning*, 21(2), 230–249.
- Bonnefon, J.-F., Feeney, A., & De Neys, W. (2011). The risk of polite misunderstandings. *Current Directions in Psychological Science*, 20(5), 321–324.
- Bonnefon, J.-F., Feeney, A., & Villejoubert, G. (2009). When some is actually all: Scalar inferences in face-threatening contexts. *Cognition*, 112(2), 249–258.
- Boyer, A. (1702). *The english theophrastus: Or, the manners of the age: Being the modern characters of the court, the town, and the city...* W. Turner... R. Basset...; J. Chantry.

- Braginsky, M., Tessler, M. H., & Hawkins, R. (n.d.). *Rwebppl: R interface to webppl*. Retrieved from <https://github.com/mhtess/rwebppl>
- Braginsky, M., Yurovsky, D., & Frank, M. C. (n.d.). *Langcog: Language and cognition lab things*. Retrieved from <http://github.com/langcog/langcog>
- Breheny, R., Katsos, N., & Williams, J. (2006). Are generalised scalar implicatures generated by default? An on-line investigation into the role of context in generating pragmatic inferences. *Cognition*, 100(3), 434–463.
- Brown, P., & Levinson, S. C. (1987). *Politeness: Some universals in language usage* (Vol. 4). Cambridge university press.
- Brown, R., & Gilman, A. (1989). Politeness theory and shakespeare’s four major tragedies. *Language in Society*, 18(2), 159–212.
- Bussey, K. (1999). Children’s categorization and evaluation of different types of lies and truths. *Child Development*, 70(6), 1338–1347.
- BÄijhler, K. (1934). *Sprachtheorie*. Oxford, England: Fischer.
- BÄijrkner, P.-C. (2017). brms: An R package for bayesian multilevel models using Stan. *Journal of Statistical Software*, 80(1), 1–28. <http://doi.org/10.18637/jss.v080.i01>
- Carlson, S. M., & Moses, L. J. (2001). Individual differences in inhibitory control and children’s theory of mind. *Child Development*, 72(4), 1032–1053.
- Cartmill, E. A., Armstrong, B. F., Gleitman, L. R., Goldin-Meadow, S., Medina, T. N., & Trueswell, J. C. (2013). Quality of early parent input predicts child vocabulary 3 years later. *Proceedings of the National Academy of Sciences*, 110(28), 11278–11283.
- Clark, H. H., & Schunk, D. H. (1980). Polite responses to polite requests. *Cognition*, 8(2), 111–143.
- Corsaro, W. A. (1979). Young children’s conception of status and role. *Sociology of Education*, 46–59.
- Culpeper, J. (2011). *Impoliteness: Using language to cause offence* (Vol. 28). Cambridge University

Press.

Dorai-Raj, S. (2014). *Binom: Binomial confidence intervals for several parameterizations*. Retrieved from <https://CRAN.R-project.org/package=binom>

Eddelbuettel, D., & Balamuta, J. J. (2017). Extending extitR with extitC++: A Brief Introduction to extitRcpp. *PeerJ Preprints*, 5, e3188v1. <http://doi.org/10.7287/peerj.preprints.3188v1>

Eddelbuettel, D., & François, R. (2011). Rcpp: Seamless R and C++ integration. *Journal of Statistical Software*, 40(8), 1–18. <http://doi.org/10.18637/jss.v040.i08>

Ervin-Tripp, S. M. (1967). An issei learns english. *Journal of Social Issues*, 23(2), 78–90.

Ervin-Tripp, S. M. (1969). Sociolinguistics. *Advances in Experimental Social Psychology*, 4, 91–165.

Ervin-Tripp, S. M. (1977). Wait for me, roller skate. In S. Ervin-Tripp & C. Mitchell-Kernan (Eds.), *Child discourse* (pp. 165–188). New York: Academic Press.

Ervin-Tripp, S. M. (1982). Ask and it shall be given unto you: Children's requests. *Georgetown University Roundtable on Languages and Linguistics. Contemporary Perceptions of Language: Interdisciplinary Dimensions*, 235–245.

Ervin-Tripp, S. M., Guo, J., & Lampert, M. (1990). Politeness and persuasion in children's control acts. *Journal of Pragmatics*, 14(2), 307–331.

Ervin-Tripp, S. M., O'Connor, M. C., & Rosenberg, J. (1984). Language and power in the family. *Language and Power*, 116–135.

Feeney, A., & Bonnefon, J.-F. (2013). Politeness and honesty contribute additively to the interpretation of scalar expressions. *Journal of Language and Social Psychology*, 32(2), 181–190.

Finley, G. E., & Humphreys, C. A. (1974). Naive psychology and the development of persuasive appeals in girls. *Canadian Journal of Behavioural Science/Revue Canadienne Des Sciences Du Comportement*, 6(1), 75.

Fisher, C. (2002). The role of abstract syntactic knowledge in language acquisition: A reply to

- tomasello (2000). *Cognition*, 82(3), 259–278.
- Franchak, J. M., Kretch, K. S., & Adolph, K. E. (2018). See and be seen: Infant–caregiver social looking during locomotor free play. *Developmental Science*, 21(4), e12626.
- Frank, M. C., & Goodman, N. D. (2012). Predicting pragmatic reasoning in language games. *Science*, 336(6084), 998–998.
- Franke, M., & Jdger, G. (2016). Probabilistic pragmatics, or why bayes’ rule is probably important for pragmatics. *Zeitschrift Fr Sprachwissenschaft*, 35(1), 3–44.
- Fraser, B. (1990). Perspectives on politeness. *Journal of Pragmatics*, 14(2), 219–236.
- Fraser, B., & Nolen, W. (1981). The association of deference with linguistic form. *International Journal of the Sociology of Language*, 1981(27), 93–110.
- Fu, G., Heyman, G. D., Chen, G., Liu, P., & Lee, K. (2015). Children trust people who lie to benefit others. *Journal of Experimental Child Psychology*, 129, 127–139.
- Gelman, A., & Hill, J. (2006). *Data analysis using regression and multilevel/hierarchical models*. Cambridge university press.
- Gleason, J. B., Perlmann, R. Y., & Greif, E. B. (1984). What’s the magic word: Learning language through politeness routines? *Discourse Processes*, 7(4), 493–502.
- Goffman, E. (1967). *Interaction ritual: Essays on face-to-face interaction*. Aldine.
- Goodman, N. D., & Frank, M. C. (2016a). Pragmatic language interpretation as probabilistic inference. *Trends in Cognitive Sciences*, 20(11), 818–829.
- Goodman, N. D., & Frank, M. C. (2016b). Pragmatic language interpretation as probabilistic inference. *Trends in Cognitive Sciences*, 20(11), 818–829.
- Goodman, N. D., & Stuhlmjller, A. (2013). Knowledge and implicature: Modeling language understanding as social cognition. *Topics in Cognitive Science*, 5(1), 173–184.
- Goodman, N. D., & Stuhlmjller, A. (2014). The Design and Implementation of Probabilistic Programming Languages. <http://dippl.org>.

- Grice, H. P. (1975). Logic and conversation. In P. Cole & J. L. Morgan (Eds.), *Syntax and semantics* (Vol. 3, pp. 41–58). Academic Press.
- Gu, Y. (1990). Politeness phenomena in modern chinese. *Journal of Pragmatics*, 14(2), 237–257.
- Halliday, M. A. K. (1975). *Learning how to mean: Explorations in the development of language*. London: Edward Arnold.
- Hamlin, K. J., Ullman, T. D., Tenenbaum, J. B., Goodman, N. D., & Baker, C. L. (2013). The mentalistic basis of core social cognition: Experiments in preverbal infants and a computational model. *Developmental Science*, 16(2), 209–226.
- Henry, L., & Wickham, H. (2017). *Purrr: Functional programming tools*. Retrieved from <https://CRAN.R-project.org/package=purrr>
- Heyman, G. D., Sweet, M. A., & Lee, K. (2009). Children’s reasoning about lie-telling and truth-telling in politeness contexts. *Social Development*, 18(3), 728–746.
- Hills, T. T., Jones, M. N., & Todd, P. M. (2012). Optimal foraging in semantic memory. *Psychological Review*, 119(2), 431.
- Hirschberg, J. B. (1985). *A theory of scalar implicature*. University of Pennsylvania.
- Hocking, T. D. (2017). *Directlabels: Direct labels for multicolor plots*. Retrieved from <https://CRAN.R-project.org/package=directlabels>
- Holtgraves, T. (1986). Language structure in social interaction: Perceptions of direct and indirect speech acts and interactants who use them. *Journal of Personality and Social Psychology*, 51(2), 305.
- Holtgraves, T. (1997). YES, but... positive politeness in conversation arguments. *Journal of Language and Social Psychology*, 16(2), 222–239.
- Holtgraves, T. (1998). Interpreting indirect replies. *Cognitive Psychology*, 37(1), 1–27.
- Holtgraves, T., & Perdeu, A. (2016). Politeness and the communication of uncertainty. *Cognition*, 154, 1–10.

- Holtgraves, T., & Yang, J.-N. (1992). Interpersonal underpinnings of request strategies: General principles and differences due to culture and gender. *Journal of Personality and Social Psychology*, 62(2), 246.
- Huang, Y. T., & Snedeker, J. (2009). Online interpretation of scalar quantifiers: Insight into the semantics–pragmatics interface. *Cognitive Psychology*, 58(3), 376–415. <http://doi.org/10.1016/j.cogpsych.2008.09.001>
- Ide, S. (1989). Formal forms and discernment: Two neglected aspects of universals of linguistic politeness. *Multilingua-Journal of Cross-Cultural and Interlanguage Communication*, 8(2-3), 223–248.
- Jakobson, R. (1960). Linguistics and poetics. In *Style in language* (pp. 350–377). MA: MIT Press.
- James, S. L. (1978). Effect of listener age and situation on the politeness of children’s directives. *Journal of Psycholinguistic Research*, 7(4), 307–317.
- Jara-Ettinger, J., Gweon, H., Schulz, L. E., & Tenenbaum, J. B. (2016). The naïve utility calculus: Computational principles underlying commonsense psychology. *Trends in Cognitive Sciences*, 20(8), 589–604.
- Kant, I. (1949). On a supposed right to lie from altruistic motives. *Critical of Practical Reason and Other Writings*, 346–350.
- Kao, J. T., & Goodman, N. D. (2015). Let’s talk (ironically) about the weather: Modeling verbal irony. In *Proceedings of the 37th annual conference of the Cognitive Science Society*.
- Kao, J. T., Wu, J. Y., Bergen, L., & Goodman, N. D. (2014). Nonliteral understanding of number words. *Proceedings of the National Academy of Sciences*, 111(33), 12002–12007.
- Kyratzis, A., & Guo, J. (2001). Preschool girls’ and boys’ verbal conflict strategies in the united states and china. *Research on Language and Social Interaction*, 34(1), 45–74.
- Lakoff, R. (1973). The logic of politeness; or, minding your p’s and q’s. In A. W. C. Corum T. Cedric Smith-Stark (Ed.), *Papers from the ninth regional meeting of the chicago linguistics society* (pp. 292–305). Chicago: Department of Linguistics, University of Chicago.

- Lassiter, D., & Goodman, N. D. (2017). Adjectival vagueness in a bayesian model of interpretation. *Synthese*, 194(10), 3801–3836.
- Lee, M. D., & Wagenmakers, E. J. (2014). *Bayesian cognitive modeling: A practical course*. Cambridge Univ. Press.
- Leech, G. (1983). *Principles of pragmatics*. London, New York: Longman Group Ltd.
- Leichty, G., & Applegate, J. L. (1991). Social-cognitive and situational influences on the use of face-saving persuasive strategies. *Human Communication Research*, 17(3), 451–484.
- Lim, T. S., & Bowers, J. W. (1991). Facework: Solidarity, approbation, and tact. *Human Communication Research*, 17(3), 415–450.
- Liu, S., Ullman, T. D., Tenenbaum, J. B., & Spelke, E. S. (2017). Ten-month-old infants infer the value of goals from the costs of actions. *Science*, 358(6366), 1038–1041.
- Locher, M. A., & Watts, R. J. (2005). Politeness theory and relational work. *Journal of Politeness Research*, 1(1), 9–33.
- Ma, F., Xu, F., Heyman, G. D., & Lee, K. (2011). Chinese children’s evaluations of white lies: Weighing the consequences for recipients. *Journal of Experimental Child Psychology*, 108(2), 308–321.
- Manohar, S. G., & Husain, M. (2013). Attention as foraging for information and value. *Frontiers in Human Neuroscience*, 7, 711.
- Mao, L. R. (1994). Beyond politeness theory: “Face” revisited and renewed. *Journal of Pragmatics*, 21(5), 451–486.
- Marchman, V. A., & Fernald, A. (2008). Speed of word recognition and vocabulary knowledge in infancy predict cognitive and language outcomes in later childhood. *Developmental Science*, 11(3), F9–F16.
- Matsumoto, Y. (1988). Reexamination of the universality of face: Politeness phenomena in japanese. *Journal of Pragmatics*, 12(4), 403–426.

- Morey, R. D., & Rouder, J. N. (2015). *BayesFactor: Computation of bayes factors for common designs*. Retrieved from <https://CRAN.R-project.org/package=BayesFactor>
- MÄijller, K. (2017a). *Bindrcpp: An 'rcpp' interface to active bindings*. Retrieved from <https://CRAN.R-project.org/package=bindrcpp>
- MÄijller, K. (2017b). *Here: A simpler way to find your files*. Retrieved from <https://CRAN.R-project.org/package=here>
- MÄijller, K., & Wickham, H. (2017). *Tibble: Simple data frames*. Retrieved from <https://CRAN.R-project.org/package=tibble>
- Neuwirth, E. (2014). *RColorBrewer: ColorBrewer palettes*. Retrieved from <https://CRAN.R-project.org/package=RColorBrewer>
- Nippold, M. A., Leonard, L. B., & Anastopoulos, A. (1982). Development in the use and understanding of polite forms in children. *Journal of Speech, Language, and Hearing Research*, 25(2), 193–202.
- Ooms, J. (2014). The jsonlite package: A practical and consistent mapping between json data and r objects. *arXiv:1403.2805 [Stat.CO]*. Retrieved from <https://arxiv.org/abs/1403.2805>
- Oudeyer, P.-Y., & Smith, L. B. (2016). How evolution may work through curiosity-driven developmental process. *Topics in Cognitive Science*, 8(2), 492–502.
- Pighin, S., & Bonnefon, J.-F. (2011). Facework and uncertain reasoning in health communication. *Patient Education and Counseling*, 85(2), 169–172.
- Pinker, S., Nowak, M. A., & Lee, J. J. (2008). The logic of indirect speech. *Proceedings of the National Academy of Sciences*, 105(3), 833–838.
- Pirolli, P., & Card, S. (1999). Information foraging. *Psychological Review*, 106(4), 643.
- Plummer, M., Best, N., Cowles, K., & Vines, K. (2006). CODA: Convergence diagnosis and output analysis for mcmc. *R News*, 6(1), 7–11. Retrieved from <https://journal.r-project.org/archive/>

- Polite. (2017a). In *OED online*. Oxford University Press. Retrieved from <http://www.oed.com/view/Entry/146878?rskey=4vSu4F&result=1&isAdvanced=false>
- Polite. (2017b). In *Cambridge online dictionary*. Cambridge University Press. Retrieved from <http://dictionary.cambridge.org/us/dictionary/english/polite>
- Quinley, J. (2011). Politeness and trust games. *Student Papers Session, Proceedings of ESSLLI*.
- R Core Team. (2017). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <https://www.R-project.org/>
- Read, B. K., & Cherry, L. J. (1978). Preschool children's production of directive forms. *Discourse Processes*, 1(3), 233–245.
- Ryckebusch, C., & Marcos, H. (2004). Speech acts, social context and parent-toddler play between the ages of 1; 5 and 2; 3. *Journal of Pragmatics*, 36(5), 883–897.
- Searle, J. (1975). Indirect speech acts. In P. Cole & J. L. Morgan (Eds.), *Syntax and semantics* (Vol. 3, pp. 59–82). Academic Press.
- Shannon, C. E. (1948). A mathematical theory of communication. *Bell Syst. Tech. J.*, 27, 623–656.
- Shatz, M., & Gelman, R. (1973). The development of communication skills: Modifications in the speech of young children as a function of listener. *Monographs of the Society for Research in Child Development*, 1–38.
- Snow, C. E., Perlmann, R. Y., Gleason, J. B., & Hooshyar, N. (1990). Developmental perspectives on politeness:: Sources of children's knowledge. *Journal of Pragmatics*, 14(2), 289–305.
- Song, M., & Song, H.-J. (2014). Five- to six-year-old children's understanding of lies and truths in different contexts. *The Korean Journal of Developmental Psychology*, 27(1), 55–71.
- Spencer-Oatey, H. (2000). *Culturally speaking: Managing rapport through talk across cultures*. London; New York: Continuum.
- Sweetser, E. (1987). The definition of lie. *Cultural Models in Language and Thought*, 43–66.
- Talwar, V., & Lee, K. (2002). Development of lying to conceal a transgression: Children's control of

- expressive behaviour during verbal deception. *International Journal of Behavioral Development*, 26(5), 436–444.
- Talwar, V., Murphy, S. M., & Lee, K. (2007). White lie-telling in children for politeness purposes. *International Journal of Behavioral Development*, 31(1), 1–11.
- Tanenhaus, M. K., Magnuson, J. S., Dahan, D., & Chambers, C. (2000). Eye movements and lexical access in spoken-language comprehension: Evaluating a linking hypothesis between fixations and linguistic processing. *Journal of Psycholinguistic Research*, 29(6), 557–580.
- Terkourafi, M. (2002). Politeness and formulaicity: Evidence from cypriot greek. *Journal of Greek Linguistics*, 3(1), 179–201.
- Turiel, E. (1977). Distinct conceptual and developmental domains: Social convention and morality. In *Nebraska symposium on motivation*. University of Nebraska Press.
- Van Rooy, R. (2003). Being polite is a handicap: Towards a game theoretical analysis of polite linguistic behavior. In *Proceedings of the 9th conference on theoretical aspects of rationality and knowledge* (pp. 45–58). ACM.
- Walper, S., & Valtin, R. (1992). Children’s understanding of white lies. In *Politeness in language: Studies in its history, theory and practice* (ed. by R. J. Watts, S. Ide, & K. Ehlich) (pp. 231–251). Mouton de Gruyter.
- Watts, R. J. (1989). Relevance and relational work: Linguistic politeness as politic behavior. *Multilingua-Journal of Cross-Cultural and Interlanguage Communication*, 8(2-3), 131–166.
- Watts, R. J. (2003). *Politeness*. Cambridge University Press.
- Watts, R. J., Ide, S., & Ehlich, K. (1992). Introduction. In *Politeness in language: Studies in its history, theory and practice* (ed. by R. J. Watts, S. Ide, & K. Ehlich) (pp. 1–17). Mouton de Gruyter.
- Wickham, H. (2009). *Ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York. Retrieved from <http://ggplot2.org>
- Wickham, H. (2017a). *Forcats: Tools for working with categorical variables (factors)*. Retrieved

- from <https://CRAN.R-project.org/package=forcats>
- Wickham, H. (2017b). *Stringr: Simple, consistent wrappers for common string operations*. Retrieved from <https://CRAN.R-project.org/package=stringr>
- Wickham, H. (2017c). *Tidymodels: Easily install and load the 'tidymodels'*. Retrieved from <https://CRAN.R-project.org/package=tidymodels>
- Wickham, H., & Henry, L. (2017). *Tidyr: Easily tidy data with 'spread()' and 'gather()' functions*. Retrieved from <https://CRAN.R-project.org/package=tidyr>
- Wickham, H., Francois, R., Henry, L., & Müller, K. (2017). *Dplyr: A grammar of data manipulation*. Retrieved from <https://CRAN.R-project.org/package=dplyr>
- Wickham, H., Hester, J., & Francois, R. (2017). *Readr: Read rectangular text data*. Retrieved from <https://CRAN.R-project.org/package=readr>
- Wilson, G., & Wood, K. (2004). The influence of children on parental purchases during supermarket shopping. *International Journal of Consumer Studies*, 28(4), 329–336.
- Xu, F., Bao, X., Fu, G., Talwar, V., & Lee, K. (2010). Lying and truth-telling in children: From concept to action. *Child Development*, 81(2), 581–596.
- Xu, F., Evans, A. D., Li, C., Li, Q., Heyman, G., & Lee, K. (2013). The role of honesty and benevolence in children's judgments of trustworthiness. *International Journal of Behavioral Development*, 37(3), 257–265.
- Yoon, E. J., MacDonald, K., Asaba, M., Gweon, H., & Frank, M. C. (2018). Balancing informational and social goals in active learning. In *Proceedings of the 40th annual conference of the cognitive science society*.
- Yoon, E. J., Tessler, M. H., Goodman, N. D., & Frank, M. C. (2017). "I won't lie, it wasn't amazing": Modeling polite indirect speech. In *Proceedings of the thirty-ninth annual conference of the Cognitive Science Society*.
- Yu, K. (2011). Culture-specific concepts of politeness: Indirectness and politeness in english, hebrew and korean requests. *Intercultural Pragmatics*, 8(3), 385–409.

- Zangl, R., & Fernald, A. (2007). Increasing flexibility in children's online processing of grammatical and nonce determiners in fluent speech. *Language Learning and Development*, 3(3), 199–231.